Optimal Carbon Tax with Dynamic Skill Heterogeneity

Alexandre Pavlov*

October 19, 2026†
Work in Progress: Latest version

Abstract

Extending a dynamic Mirrlees model to two goods, one polluting, this paper studies optimal carbon tax design when the government cares both about correcting externalities and redistributing income. Under weakly separable preferences, I show that dirty good wedges should be nonlinear and age-dependent when the dirty good affects work incentives, with their level increasing when pollution also discourages work. I also derive the optimal linear dirty good wedge, as well as the optimal dirty good wedge when labor and savings wedges are fixed at an exogenous level. In a model calibrated to US data, I however find that the best linear dirty good wedge captures 99% of the welfare gains of a nonlinear dirty good wedge. When income wedges are fixed at levels implied by the US tax system, the dirty good wedge gains more of a redistributive role and going from a linear to a nonlinear dirty good wedge increases welfare gains by 57%, of which about 74% is captured by age-dependence alone. This welfare gain is small compared to that of even modest income tax reforms however, suggesting that a more flexible carbon pricing instrument is a poor substitute for income tax reform. If income tax reform is unavailable, age-dependent carbon taxes are a good alternative to full nonlinearity, which can be hard to implement.

Keywords: *Carbon tax, dynamic Mirrlees taxation, second-best taxation*
JEL Classification: *D62; D82; H21; H23*

*Université de Montréal: alexandre.pavlov@umontreal.ca

†I thank my advisor Guillaume Sublet for his invaluable guidance and support, as well as the members of his workshop and the participants of the CIREQ Weekly PhD Student Research Workshop, the McGill University microbreakfast, the University of Montréal macro brownbag, the CIREQ-CIRANO-RRECQ Environmental and Natural Resource Economics Workshop, the 19th CIREQ PhD Students' Conference, the 59th CEA Annual Conference, the 35th CREEA Annual Conference, the 65th SCSE Conference, the 60th CEA Annual Conference, the 25th LAGV conference, the 7th WCERE conference, the 13th FAERE Annual Conference, and the 36th CREEA Annual Conference for their comments and feedback. I am grateful for the support and funding from the Fonds de Recherche du Québec Société et Culture (FRQSC) through the Doctoral Research Scholarships program. DOI: https://doi.org/10.69777/329014.

1 Introduction

Should we expect a well-paid lawyer to face the same carbon tax rate as her low-wage secretary? If the ability-to-pay principle justifies taxing the lawyer's income at a higher marginal rate, shouldn't it also call for progressive carbon taxation? Common intuition might answer yes, but a ton of carbon causes the same harm whether emitted by the rich or the poor, which suggests everyone should face the same carbon price. If consumption inequalities come from income inequalities, one could also argue that income taxation is a better way to address them. However, correction and redistribution intertwine when the composition of the consumption basket changes incentives to work and save. Hence, carbon taxation can affect the government's ability to redistribute income and income taxation can affect its ability to correct externalities. How then should the government price carbon when it cares about redistribution?

When the consumption mix does not affect labor supply decisions and income taxes are optimal, I find that a linear carbon tax at the social cost of carbon is sufficient. In the more general case where these conditions do not hold, I characterize a nonlinear age-dependent carbon tax, extending the consumption taxation results of Mirrlees (1976) and Jacobs and Mooij (2015) to a dynamic setting. When labor and savings wedges are exogenously fixed, I show how the carbon tax should adjust to partly offset these suboptimal distortions. I establish these results analytically by extending the dynamic Mirrlees model of Golosov et al. (2016) to two goods, one clean and one dirty, where the dirty good feeds a persistent pollution stock harming utility.

I then calibrate the model to US data to compare the welfare gains from relaxing restrictions on carbon and income taxation, finding that a linear carbon tax at the social cost of carbon captures nearly all the welfare gain of a nonlinear age-dependent carbon tax when income taxes are optimal. When income taxes are fixed at the calibrated schedules, nonlinear carbon taxation improves welfare over the best linear carbon tax, but this gain is small compared to the gain from optimizing income taxes while holding carbon taxes linear. Nevertheless, if carbon tax reform is the only available lever to the government, age-dependent carbon taxation captures most of the welfare gain of a fully nonlinear carbon tax. These quantitative results suggest that making carbon taxes nonlinear is a poor substitute for income tax reform, but that age-dependence can still capture most of its benefits.

My work connects to the literature on optimal income and consumption taxes. Mirrlees (1976) extends the analysis in Mirrlees (1971) to study the optimal taxation of goods, finding that relative leisure complements should be taxed more heavily to discourage their consumption and indirectly encourage work, helping redistribution. Atkinson and Stiglitz (1976) find that, when preferences are weakly separable in consumption and labor, there is no motive for differentiated consumption taxes. In a static setting with non-separable preferences, Jacobs and Boadway (2014) further show that implementing the optimal allocation generally requires nonlinear consumption taxes. Golosov et al. (2003) extend the Atkinson-Stiglitz result to a dynamic

setting with weakly separable preferences. This paper extends the Mirrlees (1976) results on consumption taxation to a dynamic setting without assuming weakly separable preferences, showing that the resulting dirty good wedge is not only nonlinear but also history-dependent.

When consumption generates pollution, consumption taxes gain an externality correction role. Cremer et al. (1998) find that the optimal marginal tax rate on the dirty good combines a corrective term and an incentive term. Jacobs and Mooij (2015) show that the social value of reducing pollution can also depend on how environmental quality interacts with labor supply. When income and consumption taxes are jointly optimal and allowed to be nonlinear, they also show that the labor wedge keeps the standard structure in Mirrlees (1971) and Saez (2001), with no additional externality correction terms. When consumption goods and environmental quality are weakly separable from labor in the utility function, their carbon tax is equal to the social cost of carbon. I extend their analysis to a dynamic setting with persistent pollution and income risk, and show that the externality introduces no additional correction terms in either the labor or the savings wedge.

These results apply to a setting where the government is free to set fully nonlinear taxes, but what happens when it cannot be so flexible? A parallel literature adopts the Ramsey approach with linear tax instruments. In a setting where the government uses linear consumption taxes to fund its expenditures, Sandmo (1975) finds that the optimal tax on the dirty good is the sum of an externality correction term and a revenue-raising Ramsey term, whereas the tax formulas for the other goods keep their structure. Bovenberg and Mooij (1994) and Bovenberg and Goulder (1996) find that the carbon tax should be lower than the social cost of carbon because it discourages labor supply, even if its revenue finances labor income tax cuts.

Barrage (2020) extends this static analysis to a dynamic setting, putting Golosov et al. (2014) in a Ramsey taxation framework. She finds an optimal formula for the carbon tax, summing production damages and utility damages divided by the marginal cost of public funds. Her calibrated model suggests the optimal carbon tax should be 8% below the first-best carbon tax. Douenne et al. (2026) and Douenne et al. (2025) extend this analysis to a heterogeneous-agent setting and find that the optimal carbon tax is nearly Pigouvian.

These results show that linear income taxation can be consistent with keeping carbon pricing close to the social cost of carbon. However, when income taxes are not set optimally, carbon taxation may help offset distortions elsewhere in the tax system. Barrage (2020) finds that, under her calibration, the optimal carbon tax is 24% below its first-best level when labor taxation is fixed and 11% below when capital taxation is fixed. Douenne et al. (2026) find that, in their calibration, fixing the labor income tax lowers the carbon tax and fixing the capital income tax raises it. In a setting with uninsurable income risk, Douenne et al. (2025) find that optimal carbon taxation remains close to the social cost of carbon even when income tax rates are fixed at observed levels. However, these studies impose a common carbon price, which may limit its ability to offset other distortions, even if there are gains to be made. I ask how much

flexibility in carbon pricing can compensate for misset labor and savings wedges, comparing the welfare gain of moving from a linear carbon tax to a nonlinear carbon tax with that of optimizing income taxes while holding carbon pricing linear.

Section 2 presents the life-cycle economy. Section 3 characterizes the optimal labor, savings, and dirty good wedges. Section 4 characterizes the dirty good wedge when labor and savings wedges are fixed. Section 5 compares the welfare gains from relaxing restrictions on carbon and income taxation, and Section 6 concludes.

2 Environment

The economy is a two-good extension of the dynamic Mirrlees setting of Golosov et al. (2016), where consumption of one of the goods contributes to a persistent pollution stock harming utility. Households face privately observed skill shocks throughout their lives and a social planner allocates consumption and labor to maximize social welfare, providing redistribution and insurance while correcting the externality. I consider a single country and a single cohort to abstract from international and intergenerational equity concerns, so time and age coincide in this model.

2.1 Households and information

There is a unit mass of households living at dates $t = 0, \ldots, T$. A household with skill $\theta_t > 0$ supplies labor l_t and earns labor income $y_t = \theta_t l_t$. Initial skill is drawn from a distribution F_0 and subsequent skills follow a Markov process $F_t(\theta|\theta_{t-1})$ with density f_t. Households are differentiated by their skill history, $\theta^t = (\theta_0, \ldots, \theta_t)$. I define Θ^t as the set of all possible skill histories up to date t.

Assumption 2.1 (Skill process)**.**
For $t = 1, \ldots, T$, the transition density $f_t(\theta|\theta_{t-1})$ is differentiable in both arguments. The following ratio is bounded as a function of θ for each previous skill θ_{t-1}, where $f_{2,t} \equiv \frac{\partial f_t}{\partial \theta_{t-1}}$.

$$\frac{\theta_{t-1} \int_\theta^\infty f_{2,t}(x|\theta_{t-1})dx}{\theta f_t(\theta|\theta_{t-1})}$$

Each household consumes a clean good c_t and a dirty good q_t, whose aggregate Q_t generates emissions contributing to a pollution stock S_t affecting welfare. Expected lifetime utility is

$$\mathbb{E}_0 \left[\sum_{t=0}^{T} \beta^t U(c_t, q_t, l_t, S_t) \right]$$

where $\beta \in (0,1)$ is the discount factor. Both consumption goods increase utility, while labor and pollution decrease it. I allow non-separability in these arguments, meaning consumption and pollution can affect each other's marginal utility as well as that of labor.

The planner observes each household's labor income and consumption of both goods, but neither skill nor labor supply[1]. A high-skill household can therefore earn as much as a low-skill household by working less, such that the planner cannot distinguish them. This creates an information problem constraining the planner's ability to provide insurance and redistribution.

Preferences satisfy a single crossing condition in labor earnings, meaning a high-skill household has a lower marginal cost of providing labor for a given consumption level and pollution stock. Let subscripts on U denote partial derivatives.

Assumption 2.2 (Preferences)**.**
U is twice continuously differentiable, with $U_c, U_q > 0$, $U_l, U_S < 0$, and $U_{cc}, U_{qq}, U_{ll}, U_{SS} \leq 0$. For given labor and pollution levels, preferences over both goods are strictly quasi-concave with

$$2U_c U_q U_{cq} - U_c^2 U_{qq} - U_q^2 U_{cc} > 0$$

Consumption of both goods satisfies Inada conditions.

$$\lim_{c\to 0} U_c(c,q,l,S) = \infty \quad \lim_{q\to 0} U_q(c,q,l,S) = \infty \quad \lim_{c\to\infty} U_c(c,q,l,S) = 0 \quad \lim_{q\to\infty} U_q(c,q,l,S) = 0$$

Single crossing in labor earnings requires that, holding arguments fixed,

$$\frac{\partial}{\partial\theta}\left[\frac{U_l(c,q,l,S)}{\theta U_c(c,q,l,S)}\right] \geq 0$$

2.2 Technology and pollution

Output, which equals labor income, can be converted into either of the consumption goods. The clean good is the numéraire and, for simplicity, both goods have price 1. A linear technology transfers resources from one period to the next at a constant gross return $R > 0$.

The clean good generates no emissions. Each unit of the dirty good generates a unit of emissions. Let F^t be the unconditional population distribution of skill histories. Aggregate emissions are

[1] Intuitively, skill may for example be viewed as talent or another unobservable characteristic. Other aspects of skill could be observable but unusable, such as physical characteristics that society might not allow the planner to act upon. Likewise, labor supply may be viewed as effort, which may be difficult to track.

$$Q_t = \int_{\Theta^{t+1}} q_t(\theta^t) dF^t(\theta^t) \qquad t = 0, \ldots, T \tag{1}$$

Current and past emissions determine the pollution stock.

$$S_t = \mathcal{S}(S_0, \ldots, S_{t-1}, Q_0, \ldots, Q_t) \tag{2}$$

$\mathcal{S}$ is differentiable and non-decreasing in each argument. Emissions can persist and affect pollution accumulation at later dates. This equation can also be rewritten as $S_t = \mathcal{S}(\mathcal{H}_t, Q_t)$, where $\mathcal{H}_t \equiv (S_0, \ldots, S_{t-1}, Q_0, \ldots, Q_{t-1})$. I do not otherwise restrict the functional form of $\mathcal{S}$. Pollution harms utility, but not output[2].

2.3 Planner's problem and wedges

Following the standard approach in the dynamic Mirrlees literature, I use a direct mechanism to characterize the allocations attainable under these information constraints. In simple terms, this essentially means that the planner asks households for their skill and then allocates consumption and earnings according to their report. So that the planner's allocation is indeed implemented, the mechanism must make truthful reporting by households optimal, meaning it must be incentive-compatible. By the revelation principle, any equilibrium allocation attainable through another mechanism, such as a tax system, can be reproduced by such a direct mechanism.

An allocation specifies quantities $c_t(\theta^t)$, $q_t(\theta^t)$, and $y_t(\theta^t)$, which are functions of the household's skill history θ^t. For given Pareto weights $\alpha(\theta) \geq 0$, the planner maximizes

$$\int_0^\infty \alpha(\theta)\, \mathbb{E}_0 \left[\sum_{t=0}^{T} \beta^t U \left(c_t(\theta^t), q_t(\theta^t), \frac{y_t(\theta^t)}{\theta_t}, S_t \right) \middle| \theta \right] dF_0(\theta) \tag{3}$$

Feasibility requires that the present value of consumption does not exceed the present value of output as well as initial resources, given by initial capital K_0.

$$\int_0^\infty \mathbb{E}_0 \left[\sum_{t=0}^{T} \frac{c_t(\theta^t) + q_t(\theta^t) - y_t(\theta^t)}{R^t} \middle| \theta \right] dF_0(\theta) \leq RK_0 \tag{4}$$

The planner's allocation must be incentive-compatible, meaning that households must report their actual skill. Let $\sigma_t(\theta^t)$ be the report of a household θ^t and let $\sigma^t(\theta^t) = (\sigma_0(\theta^0), \ldots, \sigma_t(\theta^t))$

[2]Adding output damages does not change the qualitative results of this paper. The resulting carbon tax sums an utility damage term and an output damage term, the latter unaffected by inequality.

be its reporting history up to period t. This translates to the following constraint for every household and every reporting strategy.

$$\begin{aligned}&\mathbb{E}_0\left[\sum_{t=0}^{T}\beta^t U\left(c_t(\theta^t), q_t(\theta^t), \frac{y_t(\theta^t)}{\theta_t}, S_t\right)\middle|\theta\right]\\&\geq \mathbb{E}_0\left[\sum_{t=0}^{T}\beta^t U\left(c_t(\sigma^t(\theta^t)), q_t(\sigma^t(\theta^t)), \frac{y_t(\sigma^t(\theta^t))}{\theta_t}, S_t\right)\middle|\theta\right]\end{aligned} \tag{5}$$

I focus on allocations and the marginal distortions, or wedges, they imply for the optimality conditions on labor supply, savings, and consumption decisions. However, tax implementation is beyond the scope of this paper, since it's non-unique and matching household and planner optimality conditions is insufficient to establish implementation. See Golosov et al. (2007) and Kocherlakota (2010) for a further discussion on direct mechanisms and the distinction between its allocations and tax implementation.

The planner's problem is to choose an allocation that maximizes the social welfare function (3) subject to the resource constraint (4), incentive compatibility constraints (5), and pollution accumulation constraints (2). I write $U_{c,t}$, $U_{q,t}$, and $U_{l,t}$ as the marginal utilities of the clean good, dirty good, and labor. Allocations are characterized by wedges on labor supply, savings, and the consumption mix.

$$1-\tau_t^y(\theta^t) \equiv -\frac{U_{l,t}}{\theta_t U_{c,t}} \qquad 1-\tau_t^s(\theta^t) \equiv \frac{1}{\beta R}\frac{U_{c,t}}{\mathbb{E}_t[U_{c,t+1}]} \qquad 1+\tau_t^q(\theta^t) \equiv \frac{U_{q,t}}{U_{c,t}}$$

Next section, I characterize optimal wedges as a function of elasticities and the skill process.

3 Full optimum

3.1 Recursive problem

Global incentive compatibility compares truthful reporting with every possible reporting strategy, so the constraint set is intractable. I follow the first-order approach of Kapička (2013) and Golosov et al. (2016), solving the dual problem of minimizing the present value of resource costs required to deliver a given utility with global constraints replaced by local envelope conditions.

Write $\theta^t = (\theta^{t-1}, \theta)$ and suppose that θ^{t-1} is fixed. Let $u(\theta)$ be the remaining lifetime utility after the current skill draw, where $w(\theta)$ is its continuation value.

$$u(\theta) = U\left(c(\theta), q(\theta), \frac{y(\theta)}{\theta}, S_t\right) + \beta w(\theta) \tag{6}$$

Let $\dot{u}(\theta)$ represent how remaining lifetime utility changes with the current skill draw. Incentive compatibility implies the following local envelope condition, where $w_2(\theta)$ is how the continuation value changes with current skill.

$$\dot{u}(\theta) = -U_l\left(c(\theta), q(\theta), \frac{y(\theta)}{\theta}, S_t\right)\frac{y(\theta)}{\theta^2} + \beta w_2(\theta) \tag{7}$$

These equations introduce two continuation promises that tomorrow's allocation must deliver, namely that expected utility must be $w(\theta)$ and its sensitivity with respect to current skill must be $w_2(\theta)$. The second promise is important because households with different skills can make the same report, receiving the same continuation allocation. However, under persistent shocks, they may face different future skill distributions so that a nearby household contemplating deviating towards this report may end up receiving a different expected utility.

Because the planner also made such commitments last period, its current allocation must respect the promise-keeping and threat-keeping constraints inherited from last period.

$$\hat{w}_t(\theta^{t-1}) = \int_0^\infty u(\theta) f_t(\theta|\theta_{t-1}) d\theta \tag{8}$$

$$\hat{w}_{2,t}(\theta^{t-1}) = \int_0^\infty u(\theta) f_{2,t}(\theta|\theta_{t-1}) d\theta \tag{9}$$

Let $V_t(\hat{w}_t(\cdot), \hat{w}_{2,t}(\cdot), \mathcal{H}_t)$ be the minimized resource cost of providing utility under promise-keeping and threat-keeping constraints. The planner's cost minimization problem is

$$V_t(\hat{w}_t(\cdot), \hat{w}_{2,t}(\cdot), \mathcal{H}_t) = \min_{c_t, q_t, y_t, u_t, w_t, w_{2,t}, Q_t, S_t} \int \left[c_t + q_t - y_t + \frac{1}{R} V_{t+1}(w_t(\cdot), w_{2,t}(\cdot), \mathcal{H}_{t+1})\right] dF^t$$

The planner's problem is the above subject to the remaining lifetime utility and its derivative (6)-(7), promise-keeping and threat-keeping (8)-(9), aggregate dirty good consumption (1), and pollution accumulation (2).

3.2 Labor, savings, and dirty good wedges

The externality adds no correction term to the labor and savings optimality conditions, though the addition of the second good adds an Edgeworth substitutability term to the labor wedge. Before deriving the wedges implied by the optimal allocation, I define the following elasticities.

$$\epsilon_t(\theta) = \frac{U_{ll,t} l_t}{U_{l,t}} \quad \sigma_t(\theta) = -\frac{U_{cc,t} c_t}{U_{c,t}} \quad \sigma_{cq,t}(\theta) = -\frac{U_{cq,t} q_t}{U_{c,t}} \quad \gamma_{c,t}(\theta) = \frac{U_{cl,t} l_t}{U_{c,t}} \quad \gamma_{q,t}(\theta) = \frac{U_{ql,t} l_t}{U_{q,t}}$$

The first two are the inverses of the elasticity of labor supply and of the elasticity of intertemporal substitution, the third is the level of Edgeworth substitutability between the two goods, and the last two are the complementarity between labor and each of the goods.

Proposition 3.1 (Optimal labor wedge)**.**
The labor wedge satisfies the following optimality condition

$$\frac{\tau_t^y}{1-\tau_t^y} = A_t(\theta) B_t(\theta) C_t(\theta) + \beta R \frac{\tau_{t-1}^y}{1-\tau_{t-1}^y} D_t(\theta)$$

where

$$\begin{aligned}
A_t(\theta) &= 1 + \epsilon_t(\theta) - \gamma_{c,t}(\theta) \\
B_t(\theta) &= \frac{1 - F_t(\theta)}{\theta f_t(\theta)} \\
C_t(\theta) &= \int_\theta^\infty \exp\left[\int_\theta^x \left(\sigma_t(\tilde{x}) \frac{\dot{c}_t(\tilde{x})}{c_t(\tilde{x})} + \sigma_{cq,t}(\tilde{x}) \frac{\dot{q}_t(\tilde{x})}{q_t(\tilde{x})} - \gamma_{c,t}(\tilde{x}) \frac{\dot{y}_t(\tilde{x})}{y_t(\tilde{x})}\right) d\tilde{x}\right] (1 - \lambda_{1,t} \bar{\alpha}_t(x) U_{c,t}(x)) \frac{f_t(x) dx}{1 - F_t(\theta)} \\
D_t(\theta) &= \frac{A_t(\theta)}{A_{t-1}} \frac{U_{c,t}(\theta)}{U_{c,t-1}} \frac{\theta_{t-1} \int_\theta^\infty \exp\left[-\int_\theta^x \gamma_{c,t}(\tilde{x}) \frac{d\tilde{x}}{\tilde{x}}\right] f_{2,t}(x) dx}{\theta f_t(\theta)}
\end{aligned}$$

Proof. See Appendix A.2. □

This is the standard labor wedge formula from Golosov et al. (2016), except for C_t which now includes an Edgeworth substitutability term. When dirty good consumption grows with skill, Edgeworth substitutability between the two goods leads to a higher labor wedge, all else equal. Intuitively, taking away the numéraire from households above a certain skill level makes the dirty good more enjoyable and allows the planner to extract more resources from them, as households are willing to work more to be able to consume more of the dirty good.

Proposition 3.2 (Optimal savings wedge)**.**
Let $\widehat{U}_{c,t} \equiv \frac{U_{c,t}}{1-\chi_t}$, with $\chi_t \equiv \frac{\gamma_{c,t}}{1+\epsilon_t-\gamma_{c,t}} \frac{\tau_t^y}{1-\tau_t^y}$. Savings satisfy the generalized inverse Euler condition

$$\frac{1}{\widehat{U}_{c,t}} = \frac{1}{\beta R} \mathbb{E}_t \left[\frac{1}{\widehat{U}_{c,t+1}} \right]$$

Proof. See Appendix A.3. □

This is the generalized inverse Euler condition, expressed as a reweighted inverse Euler condition similar to Hellwig (2021). When preferences are non-separable between labor and consumption, the cost of delivering utility is weighted by $1 - \chi_t$. When the numéraire is a labor complement, consumption makes work less painful and shrinks the information rent that truth-telling requires. Neither the second good nor the externality changes the structure of savings wedges.

Proposition 3.3 (Optimal dirty good wedge)**.**
The dirty good wedge between the clean and dirty good satisfies

$$\frac{\tau_t^q - \mathrm{SCC}_t}{1 + \tau_t^q} = \frac{\gamma_{c,t} - \gamma_{q,t}}{1 + \epsilon_t - \gamma_{c,t}} \frac{\tau_t^y}{1 - \tau_t^y}$$

Proof. See Appendix A.4. □

This optimality condition balances the cost of distorting consumption choices on the left-hand side against the benefit of relaxing incentive constraints on the right-hand side, after accounting for pollution. To get an intuition for this condition, consider a utility-compensated reduction in dirty good consumption and see what happens on each side.

Each unit of dirty good taken away requires $1 + \tau^q$ units of the clean good as compensation. However, it also saves a unit of this dirty good and avoids emissions, valued at the social cost of carbon. The left-hand side is therefore the marginal cost of the distortion net of the saved dirty good and avoided damages, per unit of the numéraire used as compensation.

If $\tau^y > 0$ and $\gamma_{c,t} \neq \gamma_{q,t}$, meaning there is motive for redistribution and the consumption mix affects labor disutility, the distortion also has an incentive-provision benefit. When $\gamma_{c,t} > \gamma_{q,t}$, meaning the dirty good is a relative labor substitute, a high-skill household considering to misreport as a low-skill household to be able to work less for the same income ends up valuing the dirty good more than the intended household. When the planner takes away some dirty good consumption and replaces it with more clean good consumption in this low-skill household allocation, the high-skill household may prefer to stay with its intended allocation, relaxing the incentive constraints limiting redistribution and insurance. The planner can use this to increase

the earnings it demands from households, with appropriate consumption compensation, whose cost is captured by the labor wedge on the right-hand side. $\gamma_{c,t} - \gamma_{q,t}$ measures how much changing the consumption mix discourages imitation, while $1 + \epsilon - \gamma$ measures how much increasing required earnings encourages it. The ratio reflects how much the planner can ask of households while preserving incentive compatibility. At the optimum, this marginal benefit is equalized with the marginal cost of distorting consumption choices.

When $\tau^y = 0$ or $\gamma_{c,t} = \gamma_{q,t}$, the dirty good wedge is equal to the social cost of carbon. Indeed, the right-hand side is then zero and we can rearrange the equation to get $\tau^q = \text{SCC}_t$. That is, when there is no motive for redistribution or the consumption mix does not affect labor choices, there is no reason to deviate from the standard Pigouvian rule of setting the tax on the dirty good equal to the social cost of carbon.

This is the dynamic analogue to the dirty good wedge derived in Jacobs and Mooij (2015). The intratemporal tradeoff keeps the same form, but the labor wedge term causes the right-hand side of the equation to reflect the incentive benefits on both redistribution across initial types and insurance against subsequent skill shocks. As a consequence, the dirty good wedge is not only nonlinear, but also history-dependent. This result extends the dynamic results of Golosov et al. (2003), relaxing the weak separability assumption behind their Atkinson-Stiglitz result.

3.3 Social cost of carbon

Define $\xi_t \equiv \frac{U_{Sl,t} l_t}{U_{S,t}}$ as the complementarity between environmental quality and labor supply. $\xi_t - \gamma_{c,t}$ measures this complementarity relative to that of the numéraire. When it is positive, households working more are willing to give up more clean good consumption for a marginal improvement in environmental quality. When the planner wants to redistribute income and provide insurance against income risk, the social cost of carbon values the damages caused by an additional unit of emissions, adjusted for their effect on incentives.

Proposition 3.4 (Social cost of carbon)**.**
The incentive-adjusted social cost of carbon is

$$\text{SCC}_t = -\sum_{j=0}^{T-t} \frac{1}{R^j} \mathbb{E}_0 \left[\left(1 + \frac{\xi_{t+j} - \gamma_{c,t+j}}{1 + \epsilon_{t+j} - \gamma_{c,t+j}} \frac{\tau^y_{t+j}}{1 - \tau^y_{t+j}} \right) \frac{U_{S,t+j}}{U_{c,t+j}} \frac{\partial S_{t+j}}{\partial Q_t} \right]$$

When $\tau^y_t = 0$ or $\xi_{t+j} = \gamma_{c,t+j}$ in all periods and histories, we get the standard Pigouvian rule.

$$\text{SCC}^P_t = -\sum_{j=0}^{T-t} \frac{1}{R^j} \mathbb{E}_0 \left[\frac{U_{S,t+j}}{U_{c,t+j}} \frac{\partial S_{t+j}}{\partial Q_t} \right]$$

Proof. See Appendix A.5. □

This adjusted social cost of carbon measures the direct effect of pollution on utility, as well as its indirect effect on labor provision incentives. To understand what the bracket term means, suppose the planner considers allowing a marginal increase in emissions, holding utility constant. Since each unit of pollution causes $-U_S/U_c$ in damages in units of the numéraire, the planner must compensate households with $-U_S/U_c$ more of the numéraire to keep them indifferent. Suppose environmental quality is a relative labor complement, so that $\xi - \gamma_c > 0$. Consider a high-skill household contemplating getting the allocation intended for a low-skill household. Because it would work less to produce the same income, it needs less clean consumption to compensate for the environmental harm. Therefore, the compensation it receives at its current labor supply is more than it needs to offset its valuation of the environmental harm, increasing its incentive to misreport its skill.

This effect can be offset by asking the low-skill household to generate less income and adjusting clean consumption to maintain indifference. Each unit of forgone labor income allows clean consumption to fall by $1 - \tau^y$, leaving a net resource loss of τ^y. Compensating the direct environmental harm therefore leaves an additional resource cost associated with restoring incentive compatibility. $\xi - \gamma_c$ measures how strongly pollution increases deviation incentives, while $1+\epsilon-\gamma_c$ measures how a reduction in required labor income offsets this increase. The ratio determines how demanded labor income must be adjusted to maintain incentive compatibility. When $\tau^y > 0$ and $1+\epsilon-\gamma_c > 0$, this raises the valuation of pollution above marginal damages.

When environmental quality is instead a relative labor substitute, the argument runs in reverse. This adjusted formula extends that of Jacobs and Mooij (2015) to a dynamic setting. Current emissions affect incentives not only in the current period, but also every future period when they are persistent. The social cost of carbon therefore depends on the entire life path of the labor wedge, taking into account not only intratemporal but also intertemporal redistribution.

4 Constrained optimum

The results so far let the government reach any incentive-compatible allocation, but actual carbon taxes may be set inside imperfect tax systems where the income tax is not the one a dynamic Mirrlees optimum might prescribe. To model this restriction, I consider the planner problem in which labor and savings wedges are held at exogenous schedules. That is, of the incentive-feasible allocations, the planner may only offer those producing the required wedges.

Keeping income wedges fixed rules out using carbon revenue to cut marginal income tax rates, as in the revenue-neutral swaps of British Columbia and Sweden, or to fund income-contingent transfers such as France's Chèque énergie. Since the planner chooses the best possible allocation under these restrictions, it implicitly recycles any carbon revenue optimally.

4.1 Fixed income wedges

Let the labor wedge be held at a schedule $\bar{\tau}_t^y(\theta) < 1$ with Lagrange multiplier $\eta_t(\theta) f_t(\theta|\theta_-)$.

$$\Gamma_t(\theta) \equiv -\frac{U_l\left(c(\theta), q(\theta), \frac{y(\theta)}{\theta}, S\right)}{\theta\, U_c\left(c(\theta), q(\theta), \frac{y(\theta)}{\theta}, S\right)} - (1 - \bar{\tau}_t^y(\theta)) = 0 \tag{10}$$

$\Gamma_t(\theta)$ is the distance between the labor wedge the planner would want and the fixed schedule. Moreover, the savings wedge is held at $\bar{\tau}_t^s(\theta_t) < 1$ with Lagrange multiplier $\vartheta_{t+1}(\theta^t) f_t(\theta|\theta_-)$.

$$U_{c,t}(\theta^t) = (1 - \bar{\tau}_t^s(\theta_t))\, \beta R\ \mathbb{E}_t\big[U_{c,t+1}(\theta^{t+1})\big] \tag{11}$$

Because the savings-wedge constraint involves current and future consumption, the date-t problem cannot enforce it alone. Given a marginal utility of current consumption, the constraint requires tomorrow's marginal utility to average some value M. Today's consumption is likewise constrained by the requirement $\hat{M}$ inherited from yesterday's problem. This new state is needed because two continuation plans delivering the same promised utility and marginal continuation value can still differ in their expected future marginal utility, even if the savings-wedge constraint only allows one value for the expected future marginal utility. The planner's problem is that of Section 3 with wedge constraints (10) and (11), M in the continuation value, and state $\hat{M}$.

The labor-wedge constraint adds a term to every first-order condition while the savings-wedge constraint adds two. The first term comes from $\hat{M}$ inherited from the date-$(t-1)$ problem and the second term from M given to the date-$(t+1)$ problem. These two terms can be collected into a single shadow price of the intertemporal wedge constraint:

$$\eta_t^s(\theta^t) \equiv \left[\vartheta_t - \frac{\vartheta_{t+1}(\theta^t)}{(1 - \bar{\tau}_t^s)\, \beta R^2}\right] U_{c,t}(\theta^t) \tag{12}$$

Hence, η_t prices the labor-wedge constraint and η_t^s prices the savings-wedge constraint. Both enter the first-order conditions through $\zeta_t \equiv \eta_t(1 - \bar{\tau}_t^y) - \eta_t^s$ and η_t enters once more on its own.

To simplify notation and improve interpretability, I define $\mathcal{I}_{c,t}^q = \partial/\partial c_t \ln(U_{q,t}/U_{c,t})$ as the change in the household's willingness to give up a unit of the clean good for a unit of the dirty good as it consumes more of the clean good, and $\mathcal{I}_{y,t}^q = \partial/\partial y_t \ln(U_{q,t}/U_{c,t})$ as its counterpart for labor income. $\mathcal{I}_{c,t}^S = \partial/\partial c_t \ln(-U_{S,t}/U_{c,t})$ and $\mathcal{I}_{y,t}^S = \partial/\partial y_t \ln(-U_{S,t}/U_{c,t})$ are defined similarly for environmental quality.

Proposition 4.1 (Constrained-optimal carbon price)**.**
Define $\tau_t^{y,*}$ as the desired labor wedge. The constrained-optimal dirty good wedge satisfies

$$\frac{\tau_t^q - \mathrm{SCC}_t}{1+\tau_t^q} = \frac{\gamma_{c,t} - \gamma_{q,t}}{1+\epsilon_t - \gamma_{c,t}} \frac{\tau_t^{y,*}}{1-\tau_t^{y,*}} - \eta_t \mathcal{I}_{y,t}^q - \zeta_t \mathcal{I}_{c,t}^q$$

The social cost of carbon keeps the usual form, with the fiscal adjustment term replaced by

$$1 + \frac{\xi_{t+j} - \gamma_{c,t+j}}{1+\epsilon_{t+j} - \gamma_{c,t+j}} \frac{\tau_{t+j}^{y,*}}{1-\tau_{t+j}^{y,*}} + \eta_{t+j} \mathcal{I}_{y,t+j}^S + \zeta_{t+j} \mathcal{I}_{c,t+j}^S$$

Proof. See Appendix B.1. □

The first term of the dirty good wedge has the same form as that in the full optimum, with the labor wedge being the one the planner wants but cannot necessarily achieve.

The second term reflects the shadow cost of the labor-wedge constraint. A positive η_t means that lowering the labor wedge would reduce the planner's resource cost of delivering an allocation; the planner would like to lower the labor wedge but cannot. For brevity, I will call this "overdistortion". Likewise, I will call the opposite case "underdistortion". $\mathcal{I}_{y,t}^q$ measures how willingness to pay for the dirty good changes with labor income. Suppose $\eta_t > 0$ and the dirty good is the relative labor substitute, $\mathcal{I}_{y,t}^q < 0$. Then, $-\eta_t \mathcal{I}_{y,t}^q > 0$, pulling up the dirty good wedge. Intuitively, shifting expenditure from the dirty good to the clean good makes work less painful, so a higher carbon price encourages work and undoes part of the labor distortion.

The third term captures how the consumption mix affects the two wedge constraints through the marginal utility of the clean good. ζ_t is the net shadow price of a proportional change in this marginal utility, combining the labor-wedge and savings-wedge constraint multipliers. $\mathcal{I}_{c,t}^q$ measures how willingness to pay for the dirty good changes with clean-good consumption and is positive when the dirty good is normal. Set the savings multiplier aside and suppose $\eta_t > 0$, so that $\zeta_t = \eta_t(1-\bar{\tau}_t^y) > 0$. If the dirty good is normal, $\mathcal{I}_{c,t}^q > 0$, then $-\zeta_t \mathcal{I}_{c,t}^q < 0$, pulling down the dirty good wedge. Intuitively, shifting expenditure from the clean good to the dirty good raises the marginal utility of the clean good, strengthening the incentive to work. This channel counteracts the previous one here, so the net effect of the terms depends on their relative size. Adding back η_t^s to ζ_t, $\eta_t^s < 0$ reinforces the pull, while $\eta_t^s > 0$ weakens or even reverses it.

The fiscal adjustment term to the social cost of carbon also gains two terms. If $\eta_{t+j} > 0$ and having a higher labor income raises the willingness to pay for environmental quality, $\mathcal{I}_{y,t+j}^S > 0$, then the term is positive, pulling up the social cost of carbon. If we set aside the savings multiplier η_t^s in ζ_t such that $\zeta_t > 0$ and if consuming more of the clean good raises the willingness to pay for environmental quality, $\mathcal{I}_{c,t+j}^S > 0$, then the term is once again positive, further pulling

up the social cost of carbon. Hence, the planner wants to encourage work and can do so by providing more environmental quality, which high earners and high spenders value more.

Corollary 4.1 (Constrained-optimal carbon price with separable preferences)**.**
Under weak separability or $\gamma_{c,t} = \gamma_{q,t}$,

$$\frac{\tau_t^q - \mathrm{SCC}_t}{1 + \tau_t^q} = -\zeta_t \mathcal{I}_{c,t}^q$$

When the dirty good is normal, the carbon price is below (above) it where $\zeta_t > 0$ ($\zeta_t < 0$).

Proof. Proposition 4.1 at $\gamma_{c,t} = \gamma_{q,t}$. □

If labor or savings wedges are fixed, then weak separability no longer ensures that the carbon price is set at the social cost of carbon. When consumption moves the willingness to pay for the dirty good, the carbon price can be used to offset the wedge constraints.

4.1.1 Savings-wedge constraint over the life cycle

Whereas the sign of the labor-wedge multiplier is determined by whether labor is underdistorted or overdistorted at a given date and history, acting as a local "push", that of the savings-wedge multiplier pulls in opposite directions at different dates, acting as a "pivot" across the life cycle.

Lemma 4.1 (Sign of the savings multiplier)**.**
Along any history on which η_t^s is nonzero at some date, it must take both positive and negative values over the life cycle. The result holds for any labor schedule and skill process.

Proof. See Appendix B.2. □

Intuitively, consider the case where savings is always overdistorted. If we were talking about the labor wedge, this would only affect consumption-labor choices within a single date. However, the savings wedge here frontloads consumption. Because savings is made to offer a lower return, the household consumes more today and less tomorrow. Therefore, the planner would like to reduce consumption today and increase it tomorrow. Without the means to directly affect the savings wedge, the planner's second option is to affect the consumption mix, if consumption of one good affects demand for the other good. That is, if the dirty good encourages consumption of the clean good, then the planner counteracts the frontloading by distorting the dirty good more today and less tomorrow, shifting clean good consumption from today to tomorrow. This lemma shows that both signs must appear, but does not specify their order or timing.

4.2 Linear dirty good wedge

Aside from income wedge constraints, the planner may be unable to set allocations that imply different dirty good wedges across households, equivalent to observed linear carbon taxes.

Suppose the income wedges are now free, but the allocation is such that the dirty good wedge must be common across all households at a given date. Let $s_{qq,t} < 0$ denote the compensated response of dirty good demand to its price holding labor and pollution fixed. Let us also define

$$\varpi_t \equiv \frac{\bar{\tau}_t^q - \mathrm{SCC}_t}{1 + \bar{\tau}_t^q} \qquad\qquad \varpi_t^* \equiv (\gamma_{c,t} - \gamma_{q,t}) \frac{\psi U_{c,t}}{\theta f_t}$$

Proposition 4.2 (Optimal linear dirty good wedge)**.**
With income wedges free, the linear dirty good wedge satisfies

$$\varpi_t = \frac{\mathbb{E}_0[|s_{qq,t}|\ \varpi_t^*]}{\mathbb{E}_0[|s_{qq,t}|]}$$

Proof. See Appendices C.2 and C.4. □

The optimal dirty good wedge averages the desired corrections to the social cost of carbon, with greater weight on households whose dirty-good demand responds more to a compensated price change. When the desired correction exceeds the common correction, raising the price contributes to reducing the planner's resource cost. When it falls short of this correction, this pulls the price down. The optimum balances this across households. This is similar to the multiple-agent Ramsey rule of Diamond (1975). As shown in Appendix C.4, this restriction can also change the social cost of carbon. A similar rule applies when income wedges are fixed.

Corollary 4.2 (Optimal linear dirty good wedge with fixed income wedges)**.**
Define

$$\varpi_t^{**} \equiv \varpi_t^* + (\gamma_{c,t} - \gamma_{q,t}) \frac{\eta_t}{y_t} - \zeta_t \mathcal{I}_{c,t}^q$$

With labor and savings wedges fixed, the linear dirty good wedge satisfies

$$\varpi_t = \frac{\mathbb{E}_0[|s_{qq,t}|\ \varpi_t^{**}]}{\mathbb{E}_0[|s_{qq,t}|]}.$$

Proof. See Appendices C.3 and C.4. □

Under weak separability, $\varpi_t^{**} = -\zeta_t \mathcal{I}_{c,t}^q$, so the common wedge equals the social cost of carbon when $\mathbb{E}_0[|s_{qq,t}|\zeta_t \mathcal{I}_{c,t}^q] = 0$. A wedge held constant over the lifecycle must balance marginal resource costs across households as well as dates.

In the next section, I use Proposition 4.2 and Corollary 4.2 to quantify the welfare effects of constraining the dirty good wedge to be linear across households in a single date (age-dependent carbon tax) and also across ages (linear carbon tax), using a model calibrated to US data.

5 Quantitative exercise

The theory suggests that a carbon tax should be nonlinear, but what are the likely welfare gains of implementing more flexible carbon pricing instruments? When income taxes are fixed, is using a nonlinear or even age-dependent carbon tax a good substitute for income tax reform when the latter is difficult to implement? To answer these questions, I solve the model calibrated to US data, with an energy composite good standing in for the dirty good and other types of consumption standing in for the clean good.

I follow the quantitative setup and resolution method of Farhi and Werning (2013). I extend their logarithmic consumption preferences to two goods, where the two goods are a CES aggregate of both goods with a weight on the energy good that depends on labor supply.

$$U(c, q, l) = \log(\mathcal{C}(c, q, l)) - \frac{l^\alpha}{\alpha}$$
$$\mathcal{C}(c, q, l) = \left[(1 - b_q)c^{\frac{\varsigma-1}{\varsigma}} + b_q e^{-\varkappa l^\alpha} q^{\frac{\varsigma-1}{\varsigma}}\right]^{\frac{\varsigma}{\varsigma-1}}$$

b_q is the weight of energy in the utility function, calibrated to match the observed share of energy in total consumption, about 7% in the 2023 Consumer Expenditure Survey (U.S. Bureau of Labor Statistics 2024). ς is the elasticity of substitution between the two goods, set to 0.4 so that the implied own-price elasticity is about -0.4. West and Williams (2007) report that gasoline is a relative leisure complement (labor substitute), so I set $\varkappa$ to be positive, at 0.25. Households work from age 25 to 64 and are retired from 65 to 84. Since retirees do not work, there is no more screening problem after age 64.

Since a single cohort in a single country cannot accurately capture the social cost of carbon, I fix climate damages at the social cost of carbon, which I set to be consistent with \$190/tCO$_2$ (U.S. Environmental Protection Agency 2023). This means that any deviation from the social cost of carbon is solely due to redistribution and not due to the evolution of climate damages over time. Because the social cost of carbon in the dirty good wedge is common across all households, it affects the level but not the shape of the optimal carbon tax. Hence, this simplification does not affect the qualitative distributional implications.

I compare carbon pricing when income wedges are optimized and when they are fixed. I calibrate the labor wedge using the formula developed in Heathcote et al. (2017) and Heathcote and Tsujiyama (2021), with level parameter $\lambda^T = 0.840$ and progressivity $p = 0.181$. Since there is no HSV-style calibration for the capital income tax, I set a constant savings wedge of 0.9%, equivalent to the average effective tax rate of 18% reported by the CBO[3].

Table 1: Calibration

Parameter or variable	Value	Source
Discount factor	$\beta = 0.95$	$\beta R = 1$, from Farhi and Werning (2013)
Labor-disutility curvature	$\alpha = 3$	Farhi and Werning (2013)
Log-skill innovation variance	0.0095	Farhi and Werning (2013)
Energy expenditure share	7%	U.S. Bureau of Labor Statistics (2024)
Substitution elasticity	$\varsigma = 0.4$	Labandeira et al. (2017)
Leisure complementarity	$\varkappa = 0.25$	Sign motivated by West and Williams (2007)
Labor progressivity	$p = 0.181$	Heathcote et al. (2017)
Labor schedule level	$\lambda^T = 0.840$	Heathcote and Tsujiyama (2021)
Savings wedge	$\bar{\tau}^s = 0.009$	Congressional Budget Office (2014)

Under my assumptions, as shown in Figure 1, the planner's preferred labor wedge is much more progressive than the US-calibrated schedule and the planner's preferred savings wedge is decreasing in age. Since it is plausible that the observed wedges can be rationalized with a different calibration on preferences and the skill shock process, I am not making a statement about the optimality of the US tax system. Therefore, I also consider a second exercise where income wedges are fixed at schedules that would have been optimal in a world without externalities. This simulates the case of a government that had already optimized its tax system, but cannot readjust it after introducing carbon taxes. I then compare the welfare gains from using more flexible carbon tax instruments against those of an income tax reform, measured as departures from US-calibrated income wedges towards the planner's preferred wedges.

[3] Congressional Budget Office 2014

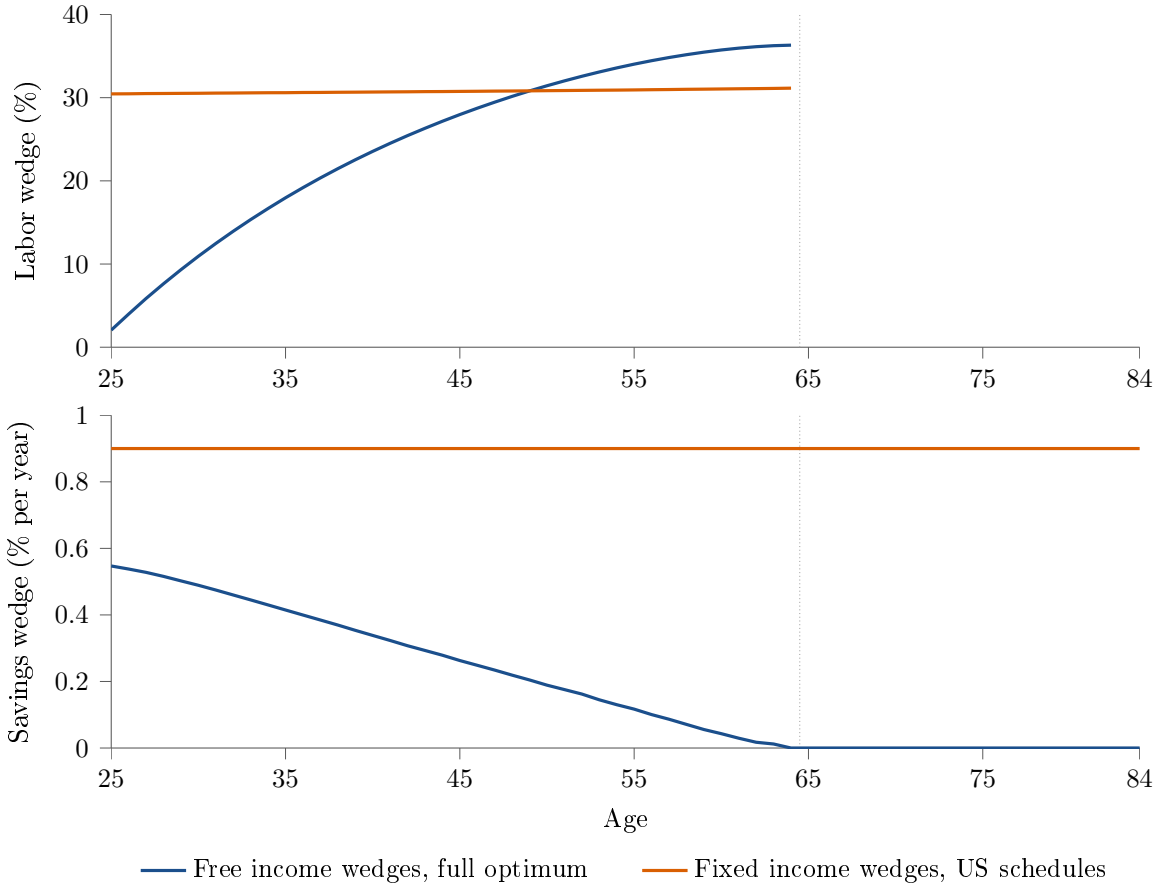


Figure 1: Average labor and savings wedges by age, with income wedges free or fixed.

5.1 Progressivity of nonlinear carbon taxes

What does an age-dependent or nonlinear history-dependent carbon tax look like and how progressive is it? Figure 2 plots the average optimal dirty good wedge over the lifecycle, comparing age-dependent and fully nonlinear dirty good wedges in a world with free labor and savings wedges and in a world with fixed labor and savings wedges.

In a world with free labor and savings wedges, the average optimal nonlinear carbon tax starts slightly above the social cost of carbon of \$190/tCO$_2$ and rises steadily to about \$245/tCO$_2$ by the end of the working life. Afterwards, it drops to the social cost of carbon, since households enter retirement and there is no more labor skill screening problem. The best age-dependent carbon tax tracks the average optimal nonlinear carbon tax almost perfectly and is therefore not plotted in the figure. Hence, under an optimal allocation, incentive-provision calls for a carbon tax that is higher than the social cost of carbon during the working life.

In a world with fixed labor and savings wedges, the average nonlinear carbon tax starts much higher than the social cost of carbon, at a bit over $\$250/\text{tCO}_2$, and declines steadily until crossing the social cost of carbon at 34 and then becoming a subsidy from age 62. This pattern comes from the fixed savings wedge. Since the savings wedge is too high at every age, households consume too much early in life, and a carbon tax that falls with age makes later consumption relatively cheaper, which works like a subsidy to saving. The age-dependent carbon tax follows a similar pattern.

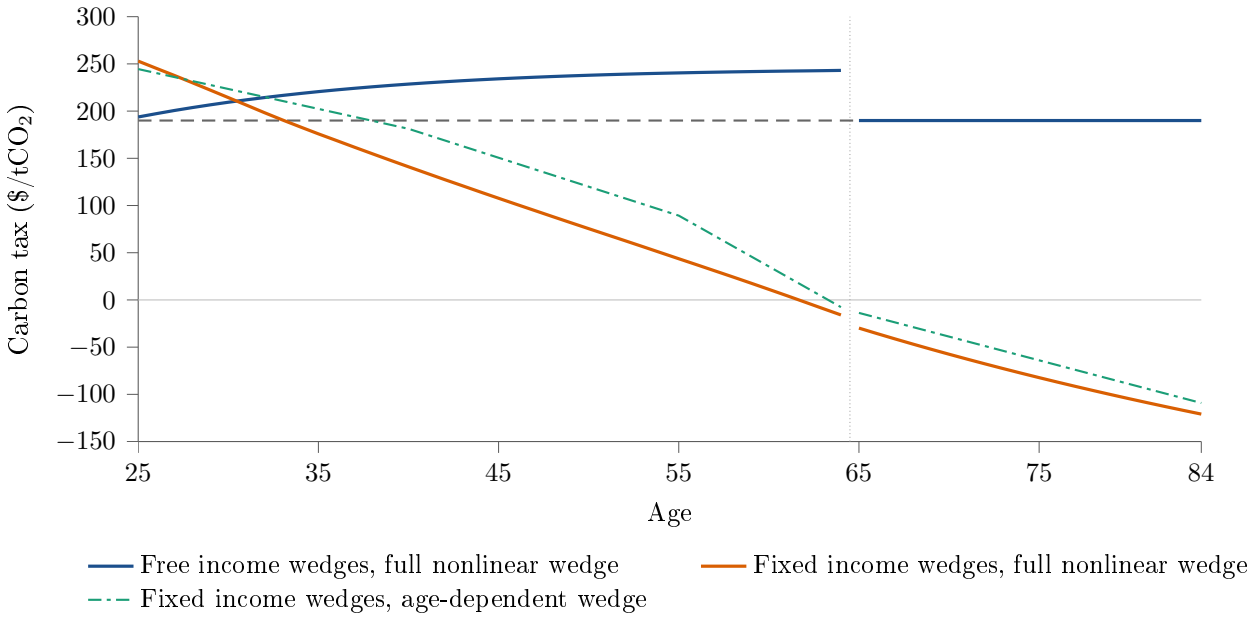


Figure 2: Average carbon tax by age. Wedges are converted to dollars. The dashed line is the social cost of carbon, $\$190/\text{tCO}_2$, and the dotted line marks retirement at age 65.

Figure 3 shows that the nonlinear carbon tax is progressive with respect to current-skill quintile, increasingly so with age. At age 64, in a world with free labor and savings wedges, the average optimal nonlinear carbon tax rises from about $\$218/\text{tCO}_2$ for the lowest-skill quintile to about $\$271/\text{tCO}_2$ for the highest-skill quintile. In a world with fixed labor and savings wedges, the average optimal nonlinear carbon tax rises from about $-\$99/\text{tCO}_2$ for the lowest-skill quintile to about $\$70/\text{tCO}_2$ for the highest-skill quintile. Hence, when income taxes are fixed, the carbon tax is about three times more progressive. While everyone pays a positive tax higher than the social cost of carbon in the first world, the first three quintiles receive a carbon subsidy rather than pay a tax in the second world.

Progressivity builds up over the working life. At age 35, the carbon tax barely varies across quintiles when income wedges are free (about \$218 to $\$222/\text{tCO}_2$), but already rises from about

$151 to $196/tCO$_2$ when they are fixed, and the lowest quintile starts receiving a subsidy on average at age 50. In short, the optimal carbon tax changes across households and ages, and may even be negative for some.

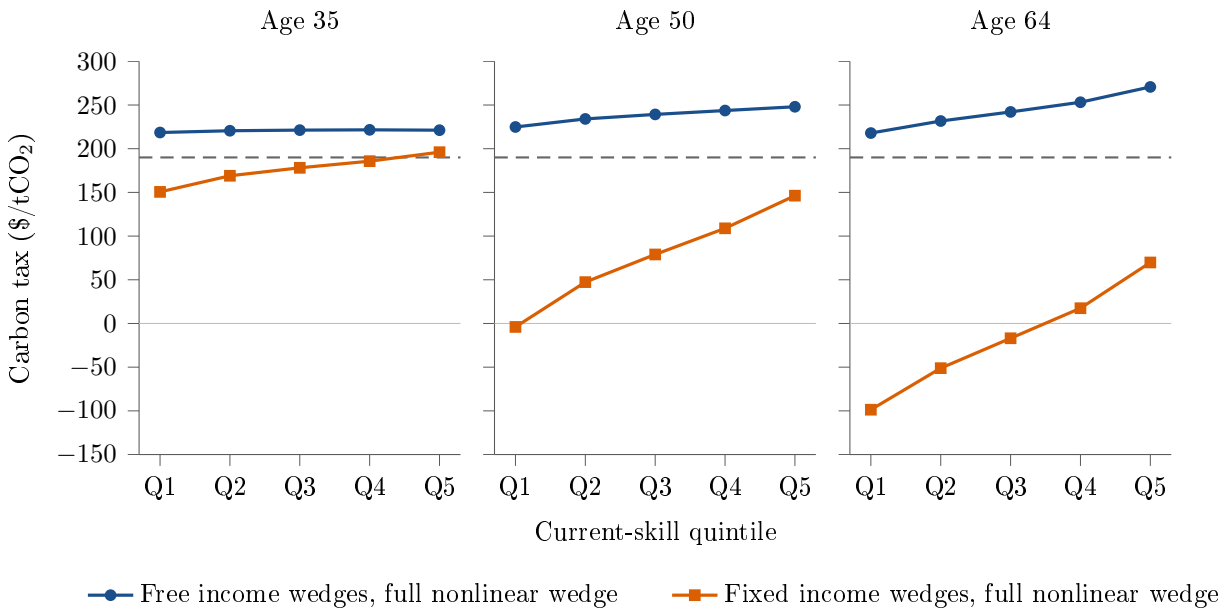


Figure 3: Average carbon tax by current-skill quintile at ages 35, 50 and 64. Wedges are converted to dollars. The dashed line is the social cost of carbon, $190/tCO$_2$.

These results suggest a significant redistributive role to carbon taxation, but this might be due to the US-calibrated income wedges being far from what the planner would like. Is there still a significant role for redistribution when the carbon tax is introduced to a world with optimized income taxes, but the latter are simply not allowed to readjust afterwards?

5.2 Dirty good wedges alongside previously optimal income wedges

Suppose the government did not use to worry about climate change and optimized its income tax system accordingly. In other words, it did its best given its information. Aftwards, it acknowledges climate change and decides to introduce a carbon tax. However, it cannot reform the entire tax system, meaning labor and savings wedges must remain fixed. Does this lack of coordination call for a change in the way the carbon tax is set?

Table 2: Gains from commodity flexibility with inherited income wedges

Carbon instrument	Externalities	No externalities
Age-dependent wedges	0.000365	0.000309
Fully nonlinear wedges	0.000559	0.000470

Note: gains in percent of lifetime consumption relative to the best linear rate.

Table 2 compares the gains from more flexible carbon instruments when labor and savings wedges are fixed at levels that would have been optimal in a world without externalities.

The best linear rate becomes about \$206/tCO_2 when income wedges can adjust, adding only 0.000059% of consumption. At the social cost of carbon, the gain from income wedges reoptimization is below 0.000001%. With income wedges frozen, age-dependent dirty good wedges range from about \$192 to \$219/tCO_2 (Figure 4). The age-dependent wedge rises during working life and falls close to the social cost of carbon at retirement. Average nonlinear wedges follow a similar path. The linear wedge, about \$205/tCO_2, captures 99.8% of the gain from introducing the optimal nonlinear wedge. Age differentiation adds 0.00037% relative to the linear wedge, about two thirds of the 0.00056% gain from the nonlinear wedge.

Having a non-zero dirty good wedge between the two goods is already useful without environmental damages. Without externalities, age differentiation adds 0.000309% relative to the best linear rate. A fully nonlinear wedge adds 0.000470% in welfare.

Hence, designing a carbon tax does not need to be paired with a redesign of the income tax system when it is already optimal. However, if it is not, is allowing for a more flexible carbon tax an adequate substitute for income tax reform?

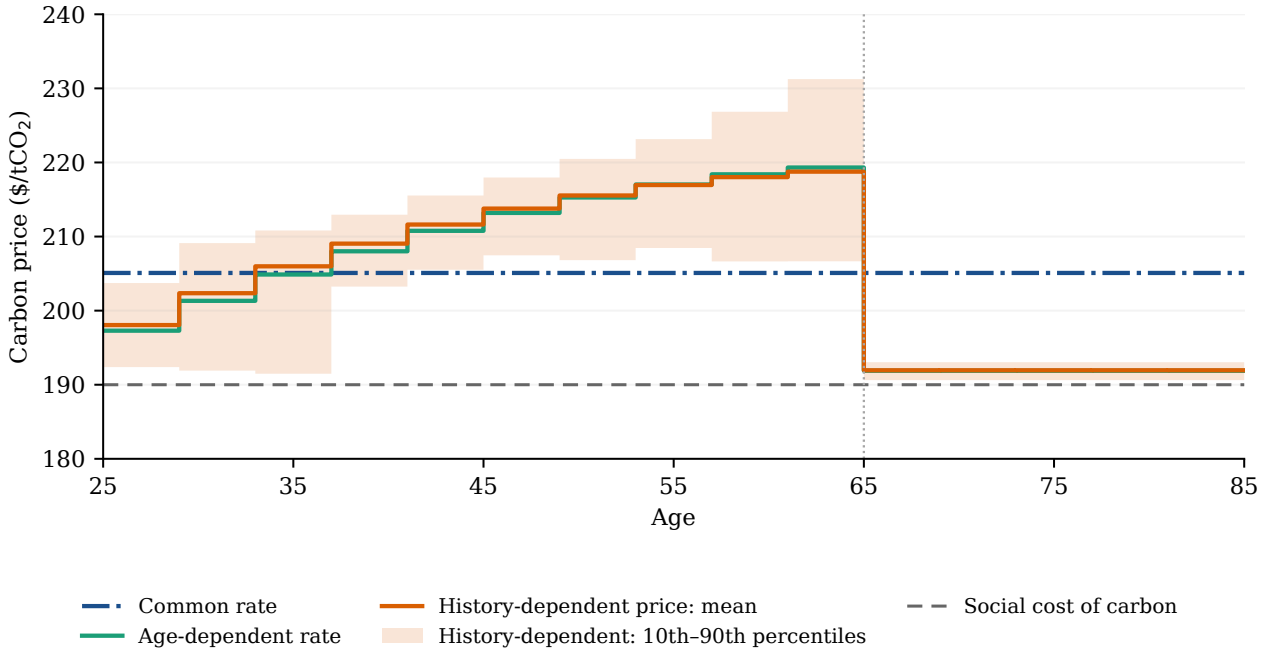


Figure 4: Optimal carbon prices with labor and savings wedges fixed at levels that would have been optimal in a world without externalities.

5.3 Income tax reform or carbon tax reform?

I come back to the US-calibrated income wedges, which the planner disagrees with. What is the welfare gain from allowing the dirty good wedge to be age-dependent, or even history-dependent? How does this compare with the welfare gain from going from the US-calibrated income wedges to the planner's preferred wedges?

Table 3 reports the welfare gains from each reform, compared against a no-tax benchmark in the US-calibrated allocation. In an optimized allocation, adding a dirty good wedge increases welfare by an appreciable 21.5% over the no-carbon benchmark. However, making the dirty good wedge age-dependent or history-dependent has little value, as the linear dirty good wedge case already captures 99% of the gain from the fully nonlinear dirty good wedge. This is because distorting consumption choices has little role in redistribution when income wedges are optimal.

On the other hand, in a US-calibrated allocation, flexibility in the dirty good wedge is more important. Allowing the dirty good wedge to be fully nonlinear increases welfare by 57.1% over the linear dirty good wedge, of which the age-dependent dirty good wedge captures about 74% of the gain. Hence, even if the fully nonlinear dirty good wedge is hard to implement, making it age-dependent gets us most of the way there.

Table 3: Welfare gains from dirty good wedge flexibility with US-calibrated income wedges and optimal wedges (% of consumption)

Income wedges	No carbon wedge	Linear carbon wedge	Age-dependent dirty good wedges	Fully nonlinear dirty good wedges
US-calibrated	0.0000	0.1913	0.2723	0.3005
Reoptimized	1.4569	1.7706	1.7720	1.7723

Note: consumption-equivalent gains relative to the US-calibrated income schedules without a carbon tax.

However, the gains from more flexible dirty good wedges are dwarfed by those of reoptimizing income wedges. Even holding dirty good wedges linear, going from the US-calibrated income wedges to the planner's preferred wedges adds 1.5793% of consumption, over 16 times the gain from making the dirty good wedge fully nonlinear. Given the large differences between the US-calibrated and planner's preferred income wedges, it might be informative to compare more modest reforms. I consider a 1% decrease in the capital income tax, a 1-2% decrease in the labor income tax, increases in the HSV progressivity parameter, capital income tax cuts in retirement since the model calls for no taxation after retirement, and graduated labor income tax cuts for younger workers since this is where the labor wedge is furthest from the planner's preferred labor wedge. Because age-dependent carbon pricing delivers most of the welfare of a fully nonlinear carbon price, I compare these gains against it. Table 4 reports the welfare gains of these limited reforms and the size of the reforms that would match the welfare gains of age-dependent carbon pricing.

The first experiment is the equivalent of a broad 1% capital income tax cut, which delivers a welfare gain of 0.2843%, exceeding the gains of age-dependent carbon pricing. In the model, this is the equivalent of reducing the savings wedge by 0.05%. The second experiment is the equivalent of a broad 1% or 2% labor income tax cut, which delivers a welfare gain of 0.2589% or 0.3200%, respectively. A cut of about 1.2% would match the welfare gain of implementing the best age-dependent carbon tax. The third experiment considers different increases in the HSV progressivity parameter, by 0.05, by 0.129, and by 0.169. The effect is smaller, reaching the welfare gains of an age-dependent carbon tax requiring increasing the parameter by about 0.132, or a 73% increase over the old parameter. The fourth experiment cuts retiree capital income taxes in half and then also considers eliminating them, with a roughly 45% cut being sufficient to match the welfare gains of an age-dependent carbon tax. The fifth experiment implements a graduated youth relief, with cuts of 2, 1.5, 1 and 0.5 points at ages 25, 26-28, 29-31 and 32-34. This particular ladder is illustrative, but is meant to take into account that the difference between US-calibrated and planner's preferred labor wedges is largest at younger ages. The resulting welfare gain is about the same as that of the age-dependent carbon tax.

Table 4: Welfare gains from limited labor wedge reforms (% of consumption)

Instrument	Reform	Gain (% of consumption)	Reform matching age-dependent carbon tax
Capital income tax	$18\% \to 17\%$	0.2843	$18\% \to 17.1\%$
Labor income tax	-1%	0.2589	-1.2%
	-2%	0.3200	
HSV progressivity (p)	$0.181 \to 0.231$	0.2240	$0.181 \to \approx 0.313$
	$0.181 \to 0.310$	0.2706	
	$0.181 \to 0.350$	0.2963	
Retiree capital income tax	$18\% \to 9\%$	0.2662	$18\% \to \approx 8\%$
	$18\% \to 0\%$	0.3285	
Graduated youth relief	From -2% at 25 to -0.5% at 32-34	0.2770	

Note: Welfare gains are measured gainst US-calibrated schedules without a carbon tax. Retirement relief leaves working-life savings wedges unchanged. The youth profile cuts labor income rates by 2, 1.5, 1 and 0.5 points at ages 25, 26-28, 29-31 and 32-34.

These examples suggest that small broad-based labor and capital tax cuts can deliver enough welfare to match the gains of an age-dependent carbon tax. If the tax system is particularly suboptimal for some ages, a targeted reform can work too. However, making the labor income tax more progressive requires an implausibly large increase in the progressivity parameter.

In short, even in a world where the income tax system implies wedges that are far from where the planner would want them to be, the welfare gain from implementing an age-dependent carbon tax is small relative to limited income tax reforms. Moreover, non-targeted income tax cuts can still be consistent with the objectives of a redistributive planner when income taxes are overall higher than where they need to be. In the US-calibrated example, labor wedges are highest for younger workers and savings wedges are always too high. The difference between observed and preferred wedges depends on the country and the problem's assumptions, but the takeaway remains that carbon pricing flexibility is a poor substitute for income tax reform.

6 Conclusion

In this paper, I explore the optimal design of carbon taxes when the government's goal is to correct for a climate externality while also redistributing income in a world where households face labor income shocks. In a dynamic Mirrlees model with two goods, I analytically show that whether households should face different carbon taxes depends on how consumption choices interact with labor supply decisions. When the dirty good is a relative labor complement, the optimal carbon tax is lower than the social cost of carbon. When labor and savings wedges are fixed at potentially suboptimal levels, the optimal carbon tax may have further correction terms offsetting these constraints.

The quantitative exercise shows the shape and progressivity of an age-dependent or fully nonlinear dirty good wedge, shows that the dirty good wedge does not need to be significantly adjusted even when income wedges were set in a world where the planner did not account for climate externalities and could not be reoptimized, and that the welfare gains from making the dirty good wedge age-dependent or history-dependent are small even relative to modest income tax reforms. Although the theory calls for nonlinear history-dependent carbon taxes, the numerical results suggest that the redistributive role of carbon taxation is limited.

Hence, at least within the confines of a Mirrlees type analysis, there is little practical benefit to deviating from the standard Pigouvian rule of setting a linear carbon tax equal to the social cost of carbon, with redistributive goals best achieved with limited income tax reform. If income taxes cannot change at all, an age-dependent carbon tax can capture most of the gains of a fully nonlinear carbon tax. In practice, this could for example be implemented by a linear carbon tax with age-dependent income-agnostic rebates.

This paper has a number of limitations that would be interesting to address in future work. First, it focuses on allocations rather than tax implementations, so it cannot speak to the specifics of how an age-dependent or even fully nonlinear carbon tax could be best put in place. Moreover, because carbon tax revenue is implicitly optimally recycled, the model cannot speak as to the best way to use it. Second, the model features a single cohort, so a natural extension is considering an overlapping generations model to study intergenerational redistribution. This would also allow for the implementation of the framework in a full integrated assessment model, in order to better study the implications of income inequality on emissions.

References

Atkinson, A.B. and J.E. Stiglitz (July 1976). "The Design of Tax Structure: Direct versus Indirect Taxation". In: *Journal of Public Economics* 6.1, pp. 55–75. ISSN: 0047-2727. DOI: 10.1016/0047-2727(76)90041-4.

Barrage, Lint (2020). "Optimal Dynamic Carbon Taxes in a Climate–Economy Model with Distortionary Fiscal Policy". In: *The Review of Economic Studies* 87.1, pp. 1–39. ISSN: 0034-6527. DOI: 10.1093/restud/rdz055.

Bovenberg, A. Lans and Ruud A. de Mooij (1994). "Environmental Levies and Distortionary Taxation". In: *The American Economic Review* 84.4. JSTOR: 2118046, pp. 1085–1089. ISSN: 00028282. (Visited on 03/14/2023).

Bovenberg, A.L. and L.H. Goulder (1996). *Optimal environmental taxation in the presence of other taxes : General equilibrium analyses*. Tech. rep. Issue: 5d4b7517-c5c8-4ef6-ab76-3. Tilburg University, School of Economics and Management. URL: https://ideas.repec.org/p/tiu/tiutis/5d4b7517-c5c8-4ef6-ab76-35964289df0c.html.

Congressional Budget Office (Dec. 2014). *Taxing Capital Income: Effective Marginal Tax Rates Under 2014 Law and Selected Policy Options*. Report. Congressional Budget Office. URL: https://www.cbo.gov/publication/49817.

Cremer, Helmuth, Firouz Gahvari, and Norbert Ladoux (Dec. 1998). "Externalities and Optimal Taxation". In: *Journal of Public Economics* 70.3, pp. 343–364. ISSN: 0047-2727. DOI: 10.1016/S0047-2727(98)00039-5.

Diamond, P.A. (Nov. 1975). "A many-person Ramsey tax rule". In: *Journal of Public Economics* 4.4, pp. 335–342. ISSN: 0047-2727. DOI: 10.1016/0047-2727(75)90009-2. URL: https://www.sciencedirect.com/science/article/pii/0047272775900092.

Douenne, Thomas, Albert Jan Hummel, and Marcelo Pedroni (Jan. 2026). "Optimal Fiscal Policy in a Climate-Economy Model with Heterogeneous Households*". In: *The Economic Journal*, ueag006. ISSN: 0013-0133. DOI: 10.1093/ej/ueag006. URL: https://doi.org/10.1093/ej/ueag006 (visited on 04/19/2026).

Douenne, Thomas et al. (2025). *Optimal Climate Policy with Incomplete Markets*. Tech. rep. Issue: tecipa-807. University of Toronto, Department of Economics. URL: https://ideas.repec.org/p/tor/tecipa/tecipa-807.html.

Farhi, Emmanuel and Iván Werning (2013). "Insurance and Taxation over the Life Cycle". In: *The Review of Economic Studies* 80.2 (283), pp. 596–635. ISSN: 00346527, 1467937X.

Golosov, Mikhail, Narayana Kocherlakota, and Aleh Tsyvinski (2003). "Optimal Indirect and Capital Taxation". In: *Review of Economic Studies* 70.3, pp. 569–587. ISSN: 0034-6527. DOI: 10.1111/1467-937x.00256.

Golosov, Mikhail, Maxim Troshkin, and Aleh Tsyvinski (2016). "Redistribution and Social Insurance". In: *American Economic Review* 106.2, pp. 359–86. DOI: 10.1257/aer.20111550.

Golosov, Mikhail, Aleh Tsyvinski, and Iván Werning (2007). "New Dynamic Public Finance: A User's Guide". In: NBER Chapters. National Bureau of Economic Research, Inc, pp. 317–388.

Golosov, Mikhail et al. (Jan. 2014). "Optimal Taxes on Fossil Fuel in General Equilibrium". In: *Econometrica* 82.1, pp. 41–88. ISSN: 0012-9682. DOI: 10.3982/ECTA10217.

Heathcote, Jonathan, Kjetil Storesletten, and Giovanni L. Violante (Nov. 2017). "Optimal Tax Progressivity: An Analytical Framework". In: *The Quarterly Journal of Economics* 132.4, pp. 1693–1754. ISSN: 0033-5533. DOI: 10.1093/qje/qjx018. URL: https://doi.org/10.1093/qje/qjx018 (visited on 07/28/2026).

Heathcote, Jonathan and Hitoshi Tsujiyama (Nov. 2021). "Optimal Income Taxation: Mirrlees Meets Ramsey". In: *Journal of Political Economy* 129.11, pp. 3141–3184. ISSN: 0022-3808. DOI: 10.1086/715851.

Hellwig, Christian (June 2021). "Static and Dynamic Mirrleesian Taxation with Non-separable Preferences: A Unified Approach". Working paper. Toulouse School of Economics.

Jacobs, Bas and Robin Boadway (Sept. 2014). "Optimal Linear Commodity Taxation under Optimal Non-Linear Income Taxation". In: *Journal of Public Economics* 117, pp. 201–210. ISSN: 0047-2727. DOI: 10.1016/j.jpubeco.2014.04.012.

Jacobs, Bas and Ruud A. de Mooij (May 2015). "Pigou Meets Mirrlees: On the Irrelevance of Tax Distortions for the Second-Best Pigouvian Tax". In: *Journal of Environmental Economics and Management* 71, pp. 90–108. ISSN: 0095-0696. DOI: 10.1016/j.jeem.2015.01.003.

Kapička, Marek (2013). "Efficient Allocations in Dynamic Private Information Economies with Persistent Shocks: A First-Order Approach". In: *The Review of Economic Studies* 80.3, pp. 1027–1054. ISSN: 0034-6527. DOI: 10.1093/restud/rds045.

Kocherlakota, Narayana R. (2010). *The New Dynamic Public Finance*. English. The Toulouse Lectures in Economics. Princeton: Princeton University Press. ISBN: 978-1-4008-3527-0. DOI: 10.1515/9781400835270. URL: https://doi.org/10.1515/9781400835270.

Labandeira, Xavier, José M. Labeaga, and Xiral López-Otero (Mar. 2017). "A meta-analysis on the price elasticity of energy demand". en. In: *Energy Policy* 102, pp. 549–568. DOI: 10.1016/j.enpol.2017.01.002. URL: https://www.sciencedirect.com/science/article/pii/S0301421517300022.

Mirrlees, J. A. (1971). "An Exploration in the Theory of Optimum Income Taxation". In: *The Review of Economic Studies* 38.2. JSTOR: 2296779, pp. 175–208. ISSN: 00346527, 1467937X. DOI: 10.2307/2296779. (Visited on 07/22/2023).

— (Nov. 1976). "Optimal Tax Theory: A Synthesis". In: *Journal of Public Economics* 6.4, pp. 327–358. ISSN: 0047-2727. DOI: 10.1016/0047-2727(76)90047-5.

Saez, Emmanuel (Jan. 2001). "Using Elasticities to Derive Optimal Income Tax Rates". In: *The Review of Economic Studies* 68.1, pp. 205–229. ISSN: 0034-6527. DOI: 10.1111/1467-937X.00166. (Visited on 07/20/2023).

Sandmo, Agnar (1975). "Optimal Taxation in the Presence of Externalities". In: *The Swedish Journal of Economics* 77.1. JSTOR: 3439329, pp. 86–98. ISSN: 00397318. DOI: 10.2307/3439329. (Visited on 02/18/2023).

U.S. Bureau of Labor Statistics (Sept. 2024). *Consumer Expenditure Surveys, 2023: Table 1101. Quintiles of Income Before Taxes: Annual Expenditure Means, Shares, Standard Errors, and Relative Standard Errors*. September 2024 release. URL: https://www.bls.gov/cex/tables/calendar-year/mean-item-share-average-standard-error/cu-income-quintiles-before-taxes-2023.xlsx.

U.S. Environmental Protection Agency (Nov. 2023). *Report on the Social Cost of Greenhouse Gases: Estimates Incorporating Recent Scientific Advances*. Report. Washington, DC: U.S. Environmental Protection Agency. URL: https://www.epa.gov/environmental-economics/scghg.

West, Sarah E. and Roberton C. Williams (Apr. 2007). "Optimal taxation and cross-price effects on labor supply: Estimates of the optimal gas tax". In: *Journal of Public Economics* 91.3, pp. 593–617. ISSN: 0047-2727. DOI: 10.1016/j.jpubeco.2006.08.007. URL: https://www.sciencedirect.com/science/article/pii/S0047272706001162.

A Optimal wedge derivations

A.1 The problem and first-order conditions

Because $U(c, q, l, S)$ can be written as $U(c, q, y/\theta, S)$, I can write U_θ as $-U_l l/\theta$. The Hamiltonian of the planner's problem of Section 3.1 is as follows, with $f_{2,t} \equiv \partial f_t / \partial \theta_{t-1}$ the persistence derivative of the transition density (Assumption 2.1), written f_2 where the date is unambiguous:

$$\begin{aligned}
\mathscr{H} = &\left[c + q - \theta l + \frac{1}{R} V_{t+1}(w, w_2, \theta, \mathbf{S})\right] f \\
&+ \psi \left[-U_l(c, q, y/\theta, S)\frac{l}{\theta} + \beta w_2\right] \\
&- \lambda_1 u(\theta) \bar{\alpha}_t f_t \\
&+ \lambda_2 u(\theta) f_2 \\
&+ \phi \left[u - U(c, q, y/\theta, S) - \beta w\right] \\
&- \mu \left[S - \mathcal{S}(\ldots, Q)\right]
\end{aligned}$$

This yields the following first-order conditions, in which $S_q \equiv \partial \mathcal{S} / \partial Q$ is the marginal effect of current aggregate emissions on the current pollution stock.

$$\begin{aligned}
\frac{\partial \mathscr{H}}{\partial c} &: f - \psi U_{cl} \frac{l}{\theta} = \phi U_c \\
\frac{\partial \mathscr{H}}{\partial q} &: f - \psi U_{ql} \frac{l}{\theta} + \left[\mu S_q + \frac{1}{R}\frac{\partial V_{t+1}}{\partial Q}\right] f = \phi U_q \\
\frac{\partial \mathscr{H}}{\partial l} &: -\theta f - \phi U_l = \frac{1}{\theta}\psi \left[U_{ll} l + U_l\right] \\
\frac{\partial \mathscr{H}}{\partial u} &: \dot{\psi} = \lambda_1 \bar{\alpha}_t f - \lambda_2 f_2 - \phi \\
\frac{\partial \mathscr{H}}{\partial w} &: \frac{1}{R}\frac{\partial V_{t+1}}{\partial w} f = \phi \beta \\
\frac{\partial \mathscr{H}}{\partial w_2} &: \frac{1}{R}\frac{\partial V_{t+1}}{\partial w_2} f = -\psi \beta \\
\frac{\partial \mathscr{H}}{\partial S} &: \int_0^\infty \left(-\psi U_{Sl}\frac{l}{\theta} - \phi U_S\right) d\theta = \mu
\end{aligned}$$

Here $S_q \equiv \partial \mathcal{S} / \partial Q$ is the marginal effect of current aggregate emissions on the current pollution stock and the density f in the first-order condition for q comes from differentiating the aggregate

$Q = \int q(\theta) f(\theta)\, d\theta$ which enters both the current accumulation constraint and, as the newest entry of the emission state $\mathbf{S}$, the continuation value through $\partial V_{t+1}/\partial Q \equiv \partial V_{t+1}/\partial Q_t$. Define the full shadow price of current emissions,

$$\nu_t \equiv \mu_t S_q + \frac{1}{R}\frac{\partial V_{t+1}}{\partial Q}. \tag{13}$$

The three wedges can be derived from here.

A.2 Labor wedge

Dividing the first-order condition for c of Appendix A.1 by U_c and using $\gamma_c \equiv U_{cl} l/U_c$, with the multipliers ϕ, ψ, λ_1, λ_2, and the boundary conditions $\psi(0) = \psi(\infty) = 0$,

$$\phi = \frac{f}{U_c} - \frac{\psi}{\theta}\gamma_c \tag{14}$$

From the first-order condition for l and since $U_{ll} l + U_l = U_l(1+\epsilon)$ with $\epsilon \equiv U_{ll} l/U_l$,

$$-\theta f - \phi U_l = \frac{\psi}{\theta} U_l (1+\epsilon) \implies \phi = -\frac{\theta f}{U_l} - \frac{\psi}{\theta}(1+\epsilon)$$

Equating the two expressions for ϕ and collecting the terms in ψ,

$$\frac{f}{U_c} + \frac{\theta f}{U_l} = -\frac{\psi}{\theta}(1+\epsilon-\gamma_c)$$

The labor wedge is $1-\tau^y = -U_l/(\theta U_c)$, so $\theta U_c/U_l = -1/(1-\tau^y)$, making the left-hand side $\frac{f}{U_c}\left(1 - \frac{1}{1-\tau^y}\right) = -\frac{f}{U_c}\frac{\tau^y}{1-\tau^y}$. Substituting and rearranging,

$$\frac{\tau^y}{1-\tau^y} = (1+\epsilon-\gamma_c)\frac{\psi U_c}{\theta f} \tag{15}$$

Next, substitute ϕ from (14) into $\dot{\psi} = \lambda_1 \bar{\alpha}_t f - \lambda_2 f_2 - \phi$,

$$\dot{\psi} - \frac{\gamma_c}{\theta}\psi = \lambda_1 \bar{\alpha}_t f - \lambda_2 f_2 - \frac{f}{U_c}$$

This is a linear first-order ordinary differential equation. Multiplying by the integrating factor $\exp\left[-\int^{\theta}\gamma_c(\tilde{x})\,d\tilde{x}/\tilde{x}\right]$ and integrating from θ to ∞ with the boundary condition $\psi(\infty)=0$,

$$\psi(\theta)=\int_{\theta}^{\infty}\exp\left[-\int_{\theta}^{x}\gamma_c(\tilde{x})\frac{d\tilde{x}}{\tilde{x}}\right]\left(\frac{f(x)}{U_c(x)}-\lambda_1\bar{\alpha}_t(x)f(x)+\lambda_2 f_2(x)\right)dx. \tag{16}$$

The remaining boundary condition $\psi(0)=0$ pins down λ_1,

$$\lambda_{1,t}=\frac{\int_0^{\infty}\exp\left[-\int_0^{x}\gamma_{c,t}(\tilde{x})\frac{d\tilde{x}}{\tilde{x}}\right]\left(\frac{1}{U_{c,t}(x)}f_t(x)+\lambda_{2,t}f_{2,t}(x)\right)dx}{\int_0^{\infty}\exp\left[-\int_0^{x}\gamma_{c,t}(\tilde{x})\frac{d\tilde{x}}{\tilde{x}}\right]\bar{\alpha}_t(x)f_t(x)dx}.$$

Substituting (16) into (15),

$$\frac{\tau^y}{1-\tau^y}=(1+\epsilon-\gamma_c)\frac{U_c(\theta)}{\theta f(\theta)}\int_{\theta}^{\infty}\exp\left[-\int_{\theta}^{x}\gamma_c(\tilde{x})\frac{d\tilde{x}}{\tilde{x}}\right]\left(\frac{f(x)}{U_c(x)}-\lambda_1\bar{\alpha}_t(x)f(x)+\lambda_2 f_2(x)\right)dx$$

Next, I substitute λ_1 above and rearrange. I note that

$$\frac{d}{d\theta}\ln U_{c,t}=\frac{U_{cc}\dot{c}+U_{cq}\dot{q}+U_{cl}\dot{l}}{U_c}=-\sigma_t\frac{\dot{c}}{c}-\sigma_{cq,t}\frac{\dot{q}}{q}+\gamma_c\frac{\dot{l}}{l}$$

Since $l=y/\theta$ gives $\dot{l}/l=\dot{y}/y-1/\theta$, integrating from θ to x and multiplying by the integrating factor gives the following:

$$\frac{U_{c,t}(\theta)}{U_{c,t}(x)}\exp\left[-\int_{\theta}^{x}\gamma_c(\tilde{x})\frac{d\tilde{x}}{\tilde{x}}\right]=\exp\left[\int_{\theta}^{x}\left(\sigma_t(\tilde{x})\frac{\dot{c}_t(\tilde{x})}{c_t(\tilde{x})}+\sigma_{cq,t}(\tilde{x})\frac{\dot{q}_t(\tilde{x})}{q_t(\tilde{x})}-\gamma_c(\tilde{x})\frac{\dot{y}_t(\tilde{x})}{y_t(\tilde{x})}\right)d\tilde{x}\right],$$

The first part of the labor wedge is therefore

$$(1+\epsilon_t(\theta)-\gamma_{c,t}(\theta))\,\frac{1}{\theta f_t(\theta)}\int_{\theta}^{\infty}\exp\left[\int_{\theta}^{x}\left(\sigma_t(\tilde{x})\frac{\dot{c}_t(\tilde{x})}{c_t(\tilde{x})}+\sigma_{cq,t}(\tilde{x})\frac{\dot{q}_t(\tilde{x})}{q_t(\tilde{x})}-\gamma_{c,t}(\tilde{x})\frac{\dot{y}_t(\tilde{x})}{y_t(\tilde{x})}\right)d\tilde{x}\right]$$
$$\left(1-\lambda_{1,t}\bar{\alpha}_t(x)U_{c,t}(x)\right)f_t(x)\,dx$$

For the second part of the labor wedge, the first-order condition for w_2 rearranges to $\partial V_{t+1}/\partial w_2 = -\beta R\,\psi/f$ and the envelope theorem applied to the period-$(t+1)$ problem gives $\partial V_{t+1}/\partial \hat{w}_2 = -\lambda_{2,t+1}$. Equating the two expressions and shifting the date back one period,

$$\lambda_{2,t}(\theta^{t-1}) = \beta R\,\frac{\psi_{t-1}}{f_{t-1}}$$

Next, evaluating (15) at $t-1$,

$$\frac{\psi_{t-1}}{f_{t-1}} = \frac{\tau^y_{t-1}}{1-\tau^y_{t-1}}\,\frac{\theta_{t-1}}{U_{c,t-1}\left(1+\epsilon_{t-1}-\gamma_{c,t-1}\right)}$$

The $\lambda_2 f_2$ term of the wedge is $\lambda_{2,t}$, a constant in the current type, times the same terms before the parenthetic term.

$$\left(1+\epsilon_t(\theta)-\gamma_{c,t}(\theta)\right)\frac{U_{c,t}(\theta)}{\theta f_t(\theta)}\,\lambda_{2,t}\int_\theta^\infty \exp\left[-\int_\theta^x \gamma_{c,t}(\tilde{x})\frac{d\tilde{x}}{\tilde{x}}\right] f_{2,t}(x)\,dx$$

Substituting the two and rearranging,

$$\beta R\,\frac{\tau^y_{t-1}}{1-\tau^y_{t-1}}\,\frac{1+\epsilon_t(\theta)-\gamma_{c,t}(\theta)}{1+\epsilon_{t-1}-\gamma_{c,t-1}}\,\frac{U_{c,t}(\theta)}{U_{c,t-1}}\,\frac{\theta_{t-1}\int_\theta^\infty \exp\left[-\int_\theta^x \gamma_{c,t}(\tilde{x})\frac{d\tilde{x}}{\tilde{x}}\right] f_{2,t}(x)\,dx}{\theta f_t(\theta)}$$

Putting together the two and multiplying and dividing the first term by $1-F_t(\theta)$,

$$\frac{\tau^y_t}{1-\tau^y_t} = A_t(\theta)B_t(\theta)C_t(\theta) + \beta R\frac{\tau^y_{t-1}}{1-\tau^y_{t-1}}D_t(\theta)$$

with

$$\begin{aligned}
A_t(\theta) &= 1+\epsilon_t(\theta)-\gamma_{c,t}(\theta)\\
B_t(\theta) &= \frac{1-F_t(\theta)}{\theta f_t(\theta)}\\
C_t(\theta) &= \int_\theta^\infty \exp\left[\int_\theta^x\left(\sigma_t(\tilde{x})\frac{\dot{c}_t(\tilde{x})}{c_t(\tilde{x})}+\sigma_{cq,t}(\tilde{x})\frac{\dot{q}_t(\tilde{x})}{q_t(\tilde{x})}-\gamma_{c,t}(\tilde{x})\frac{\dot{y}_t(\tilde{x})}{y_t(\tilde{x})}\right)d\tilde{x}\right]\left(1-\lambda_{1,t}\bar{\alpha}_t(x)U_{c,t}(x)\right)\frac{f_t(x)dx}{1-F_t(\theta)}\\
D_t(\theta) &= \frac{A_t(\theta)}{A_{t-1}}\frac{U_{c,t}(\theta)}{U_{c,t-1}}\frac{\theta_{t-1}\int_\theta^\infty \exp\left[-\int_\theta^x \gamma_{c,t}(\tilde{x})\frac{d\tilde{x}}{\tilde{x}}\right] f_{2,t}(x)dx}{\theta f_t(\theta)}
\end{aligned}$$

A.3 Savings wedge

The savings wedge follows from the same first-order conditions of Appendix A.1. I define χ_t, a term reducing the cost of delivering utility when consumption complements labor.

$$\chi_t(\theta^t) \equiv \frac{\gamma_{c,t}}{1+\epsilon_t-\gamma_{c,t}} \frac{\tau_t^y}{1-\tau_t^y}$$

Using $\frac{\tau_t^y}{1-\tau_t^y} = \frac{\psi U_c}{\theta f}(1+\epsilon_t-\gamma_{c,t})$, $\phi = f/U_c - \psi\gamma_c/\theta$ from (14) can be rewritten as follows:

$$\frac{\phi}{f} = \frac{1}{U_c}\left(1-\gamma_c\frac{\psi U_c}{\theta f}\right) = \frac{1-\chi_t}{U_{c,t}} \tag{17}$$

$\chi_t > 0$ when the numéraire is complementary with labor and labor is distorted, and zero when preferences are separable between the numéraire and labor or when labor is undistorted. The first-order condition for w and the envelope condition $\partial V_t/\partial\hat{w} = \lambda_{1,t}$ one period ahead give

$$\lambda_{1,t+1}(\theta^t) = \beta R\frac{\phi(\theta)}{f(\theta)} = \beta R\frac{1-\chi_t}{U_{c,t}}. \tag{18}$$

With $\psi(0) = \psi(\infty) = 0$,

$$0 = \int_0^\infty \dot{\psi}\, d\theta = \lambda_{1,t}\int_0^\infty \bar{\alpha}_t f\, d\theta - \lambda_{2,t}\int_0^\infty f_2\, d\theta - \int_0^\infty \phi\, d\theta$$

$\lambda_{1,t} = \int_0^\infty \phi\, d\theta$, which by (17) is

$$\lambda_{1,t} = \mathbb{E}_{t-1}\left[\frac{1-\chi_t}{U_{c,t}}\right] \tag{19}$$

$\mathbb{E}_{t-1}$ integrates over θ_t given θ_{t-1}. Evaluating (19) at $t+1$ and equating with (18) gives:

$$\frac{1-\chi_t}{U_{c,t}} = \frac{1}{\beta R}\mathbb{E}_t\left[\frac{1-\chi_{t+1}}{U_{c,t+1}}\right].$$

Where $1-\chi_t > 0$, defining the incentive-adjusted marginal utility $\widehat{U}_{c,t} \equiv U_{c,t}/(1-\chi_t)$,

$$\frac{1}{\widehat{U}_{c,t}} = \frac{1}{\beta R}\mathbb{E}_t\left[\frac{1}{\widehat{U}_{c,t+1}}\right]$$

A.4 dirty good wedge

Starting from the first-order conditions of Appendix A.1,

$$\phi = \frac{f}{U_c} - \frac{\psi U_{cl} \frac{l}{\theta}}{U_c}$$

Because c and q share the same ϕ,

$$\phi = \frac{f}{U_q} - \frac{\psi U_{ql} \frac{l}{\theta}}{U_q} + \frac{\nu_t f}{U_q}$$

Hence,

$$\frac{f}{U_q} - \frac{\psi U_{ql} \frac{l}{\theta}}{U_q} + \frac{\nu_t f}{U_q} = \frac{f}{U_c} - \frac{\psi U_{cl} \frac{l}{\theta}}{U_c}$$

Multiplying by U_q and U_c,

$$U_c \left(f - \psi U_{ql} \frac{l}{\theta} + \nu_t f \right) = U_q \left(f - \psi U_{cl} \frac{l}{\theta} \right)$$

Using $\gamma_c \equiv \frac{U_{cl}\, l}{U_c}$ and $\gamma_q \equiv \frac{U_{ql}\, l}{U_q}$ or, equivalently, $\gamma_c U_c = U_{cl}\, l$ and $\gamma_q U_q = U_{ql}\, l$,

$$U_c \left(f - \frac{\psi}{\theta} \gamma_q U_q + \nu_t f \right) = U_q \left(f - \frac{\psi}{\theta} \gamma_c U_c \right)$$

Then, putting the γ terms (as well as the externality term) on one side,

$$U_c f - U_q f = \frac{\psi}{\theta} U_c U_q (\gamma_q - \gamma_c) - \nu_t U_c f$$

Finally, dividing by $U_q f$,

$$\frac{U_c}{U_q} - 1 = \frac{\psi U_c}{\theta f} (\gamma_q - \gamma_c) - \nu_t \frac{U_c}{U_q}$$

Note that, by definition, $1 + \tau^q = U_q / U_c$, so $\frac{U_c}{U_q} - 1 = -\frac{\tau^q}{1+\tau^q}$. Multiplying through by -1,

$$\frac{\tau^q}{1+\tau^q} = \frac{\psi U_c}{\theta f} (\gamma_c - \gamma_q) + \nu_t \frac{U_c}{U_q}$$

Since $\frac{U_c}{U_q} = \frac{1}{1+\tau^q}$,

$$\frac{\tau^q}{1+\tau^q} = \frac{\psi U_c}{\theta f} (\gamma_c - \gamma_q) + \frac{\nu_t}{1+\tau^q}$$

Rearranging,

$$\frac{\tau^q - \nu_t}{1+\tau^q} = \frac{\psi U_c}{\theta f}(\gamma_c - \gamma_q)$$

Note that, from the labor wedge derivation,

$$\frac{\psi U_c}{\theta f} = \frac{1}{1+\epsilon_t - \gamma_{c,t}} \frac{\tau_t^y}{1-\tau_t^y}$$

Substituting,

$$\frac{\tau_t^q - \nu_t}{1+\tau_t^q} = \frac{\gamma_{c,t} - \gamma_{q,t}}{1+\epsilon_t - \gamma_{c,t}} \frac{\tau_t^y}{1-\tau_t^y}$$

A.5 The social cost of carbon

Now consider the shadow price of current emissions ν_t of (13). Its first piece is the date-t pollution-stock multiplier μ_t, the marginal damage of a unit of the current pollution stock. Substituting ϕ from the first-order condition for c of Appendix A.1 into the first-order condition for S, then dividing and multiplying by f,

$$\mu_t = \int_0^\infty \left(-\psi U_{Sl} \frac{l}{\theta} - \left[\frac{f}{U_c} - \frac{\psi U_{cl} \frac{l}{\theta}}{U_c} \right] U_S \right) d\theta = \int_0^\infty \left(-\frac{\psi}{\theta f} U_{Sl} l - \left[1 - \frac{\psi}{\theta f} U_{cl} l \right] \frac{U_S}{U_c} \right) f \, d\theta$$

Rewrite $\frac{\psi}{\theta f}$ using (15),

$$\frac{\psi}{\theta f} = \left(\frac{\tau^y}{1-\tau^y} \right) \frac{1}{U_c} \frac{1}{1+\epsilon-\gamma_c}$$

Hence, also defining $\xi \equiv \frac{U_{Sl} l}{U_S}$,

$$
\begin{aligned}
\mu_t &= \int \left(- \left(\frac{\tau^y}{1-\tau^y} \right) \frac{1}{1+\epsilon-\gamma_c} \frac{U_{Sl} l}{U_c} - \left[1 - \left(\frac{\tau^y}{1-\tau^y} \right) \frac{1}{1+\epsilon-\gamma_c} \frac{U_{cl} l}{U_c} \right] \frac{U_S}{U_c} \right) f \, d\theta \\
&= \int \left(- \left(\frac{\tau^y}{1-\tau^y} \right) \frac{1}{1+\epsilon-\gamma_c} \frac{U_{Sl} l}{U_c} \frac{U_S}{U_S} - \left[1 - \left(\frac{\tau^y}{1-\tau^y} \right) \frac{\gamma_c}{1+\epsilon-\gamma_c} \right] \frac{U_S}{U_c} \right) f \, d\theta \\
&= \int \left(- \left(\frac{\tau^y}{1-\tau^y} \right) \frac{\xi}{1+\epsilon-\gamma_c} \frac{U_S}{U_c} - \left[1 - \left(\frac{\tau^y}{1-\tau^y} \right) \frac{\gamma_c}{1+\epsilon-\gamma_c} \right] \frac{U_S}{U_c} \right) f \, d\theta \\
&= - \int \left[1 + \left(\frac{\tau^y}{1-\tau^y} \right) \frac{\xi-\gamma_c}{1+\epsilon-\gamma_c} \right] \frac{U_S}{U_c} f \, d\theta
\end{aligned}
$$

The integral runs over the date-t population of skill histories, meaning that pollution damages are weighted aggregate across all households.

We have $\mu S_q + \frac{1}{R} \frac{\partial V_{t+1}}{\partial Q}$, so we need to evaluate $\frac{\partial V_{t+1}}{\partial Q}$. By the envelope theorem applied to the period-$(t+1)$ problem, Q_t enters continuation costs only through the accumulation constraints of later dates, $S_{t+j} = \widetilde{\mathcal{S}}_{t+j}(Q_0, \ldots, Q_{t+j})$, which carry Q_t as an argument up to the final date. Hence,

$$
\frac{\partial V_{t+1}}{\partial Q_t} = \mu_{t+1} \frac{\partial S_{t+1}}{\partial Q_t} + \frac{1}{R} \frac{\partial V_{t+2}}{\partial Q_t}
$$

Iterating to the terminal date, where $V_{T+1} = 0$,

$$
\frac{\partial V_{t+1}}{\partial Q_t} = \sum_{j=1}^{T-t} \frac{1}{R^{j-1}} \mu_{t+j} \frac{\partial S_{t+j}}{\partial Q_t}
$$

Substituting into (13),

$$
\nu_t = \mu_t \frac{\partial S_t}{\partial Q_t} + \frac{1}{R} \sum_{j=1}^{T-t} \frac{1}{R^{j-1}} \mu_{t+j} \frac{\partial S_{t+j}}{\partial Q_t} = \sum_{j=0}^{T-t} \frac{1}{R^j} \mu_{t+j} \frac{\partial S_{t+j}}{\partial Q_t}
$$

Substituting for each μ_{t+j} and writing the integrals as expectations over skill histories,

$$
\mathrm{SCC}_t \equiv \nu_t = - \sum_{j=0}^{T-t} \frac{1}{R^j} \, \mathbb{E}_0 \left[\left(1 + \frac{\xi_{t+j} - \gamma_{c,t+j}}{1 + \epsilon_{t+j} - \gamma_{c,t+j}} \frac{\tau^y_{t+j}}{1 - \tau^y_{t+j}} \right) \frac{U_{S,t+j}}{U_{c,t+j}} \frac{\partial S_{t+j}}{\partial Q_t} \right]
$$

B Derivations with fixed income wedges

B.1 Constrained-optimal dirty good wedge

The labor-wedge constraint involves only the current allocation and pollution stock, so it adds a pointwise multiplier but no new state. For savings, define the following for $t = 0, \ldots, T-1$:

$$M(\theta^t) \equiv \frac{U_{c,t}(\theta^t)}{(1 - \bar{\tau}_t^s(\theta_t))\beta R}.$$

The savings-wedge constraint requires $\mathbb{E}_t[U_{c,t+1}] = M(\theta^t)$. Last period's problem brings $\hat{M}$, which enters as a new state.

$$\hat{M} = \int U_c \left(c(\theta), q(\theta), \frac{y(\theta)}{\theta}, S \right) f_t(\theta \mid \theta_-) \, d\theta$$

Because there is no previous period at birth and no future period at death, $\vartheta_0 = \vartheta_{T+1} = 0$. The Hamiltonian is that of Appendix A.1 with the two wedge constraints.

$$\begin{aligned}
\mathscr{H} = & \left[c + q - \theta l + \frac{1}{R} V_{t+1}(w, w_2, M(c, q, l, S), \theta, \mathbf{S}) \right] f \\
& + \psi \left[-U_l \frac{l}{\theta} + \beta w_2 \right] - \lambda_1 u \bar{\alpha}_t f + \lambda_2 u f_2 + \phi[u - U - \beta w] - \mu[S - \mathcal{S}(\ldots, Q)] \\
& + \eta \left[-\frac{U_l}{\theta U_c} - (1 - \bar{\tau}_t^y(\theta)) \right] f + \vartheta_t (U_c - \hat{M}) f.
\end{aligned}$$

For $z \in \{c, q, l, S\}$, the new constraints and the modified continuation value contribute

$$\left[\vartheta_t - \frac{\vartheta_{t+1}}{(1 - \bar{\tau}_t^s)\beta R^2} \right] U_{cz} f = \frac{\eta_t^s}{U_c} U_{cz} f$$

From here,

$$\begin{aligned}
\frac{\mathscr{H}}{\partial c}: &\quad f - \psi U_{cl}\frac{l}{\theta} + \eta f \Gamma_c + \frac{\eta_t^s}{U_c} U_{cc} f = \phi U_c \\
\frac{\mathscr{H}}{\partial q}: &\quad (1+\nu_t) f - \psi U_{ql}\frac{l}{\theta} + \eta f \Gamma_q + \frac{\eta_t^s}{U_c} U_{cq} f = \phi U_q \\
\frac{\mathscr{H}}{\partial l}: &\quad -\theta f - \phi U_l + \eta f \Gamma_l + \frac{\eta_t^s}{U_c} U_{cl} f = \frac{\psi}{\theta}(U_{ll} l + U_l), \\
\frac{\mathscr{H}}{\partial u}: &\quad \dot{\psi} = \lambda_1 \bar{\alpha}_t f - \lambda_2 f_2 - \phi \\
\frac{\mathscr{H}}{\partial w}: &\quad \frac{1}{R} V_{t+1,w} f = \phi\beta \\
\frac{\mathscr{H}}{\partial w_2}: &\quad \frac{1}{R} V_{t+1,w_2} f = -\psi\beta \\
\frac{\mathscr{H}}{\partial S}: &\quad \mathbb{E}_0\left[-\frac{\psi}{f} U_{Sl}\frac{l}{\theta} - \frac{\phi}{f} U_S + \eta \Gamma_S + \frac{\eta_t^s}{U_c} U_{cS}\right] = \mu_t
\end{aligned}$$

Here, $\nu_t = \mu_t S_q + R^{-1} V_{t+1,Q}$. The pollution condition averages component contributions over all date-t histories. At the constrained allocation, $\Gamma_z = -U_{lz}/(\theta U_c) - (1-\bar{\tau}_t^y) U_{cz}/U_c$. Hence,

$$\begin{aligned}
\Gamma_c &= (1-\bar{\tau}_t^y)\frac{\sigma}{c} - \frac{\gamma_c}{y} \\
\Gamma_q &= (1-\bar{\tau}_t^y)\frac{\sigma_{cq}}{q} - \frac{\gamma_q(1+\tau^q)}{y} \\
\Gamma_l &= (1-\bar{\tau}_t^y)\frac{\epsilon - \gamma_c}{l} \\
\Gamma_S &= -\frac{U_S}{U_c}\left[\frac{\xi}{y} + (1-\bar{\tau}_t^y)\frac{\xi_c}{c}\right]
\end{aligned}$$

Eliminating ϕ between the consumption and dirty-good conditions gives

$$\frac{\tau_t^q - \nu_t}{1+\tau_t^q} = (\gamma_{c,t} - \gamma_{q,t})\frac{\psi U_{c,t}}{\theta f_t} - \eta_t \Gamma_{c,t} + \eta_t^s \frac{\sigma_t}{c_t} + \frac{\eta_t \Gamma_{q,t} - \eta_t^s \frac{\sigma_{cq,t}}{q_t}}{1+\tau_t^q}$$

Putting the η_t and η_t^s terms together,

$$\frac{\tau_t^q - \nu_t}{1+\tau_t^q} = (\gamma_{c,t} - \gamma_{q,t})\frac{\psi U_{c,t}}{\theta f_t} - \eta_t\left[\Gamma_{c,t} - \frac{\Gamma_{q,t}}{1+\tau_t^q}\right] + \eta_t^s\left[\frac{\sigma_t}{c_t} - \frac{\sigma_{cq,t}}{q_t(1+\tau_t^q)}\right]$$

Using the definitions for Γ_c and Γ_q,

$$\Gamma_{c,t} - \Gamma_{q,t} = (1 - \bar{\tau}_t^y) \left[\frac{\sigma_{c,t}}{c_t} - \frac{\sigma_{cq,t}}{q_t(1+\tau_t^q)} \right] - \frac{\gamma_{c,t} - \gamma_{q,t}}{y_t}$$

Note that the first term in the brackets is the same as the η_t^s term. If we put them together,

$$\frac{\tau_t^q - \nu_t}{1+\tau_t^q} = (\gamma_{c,t} - \gamma_{q,t}) \frac{\psi U_{c,t}}{\theta f_t} + \eta_t \frac{\gamma_{c,t} - \gamma_{q,t}}{y_t} - ((1 - \hat{\tau}_t^y)\eta_t - \eta_t^s) \left[\frac{\sigma_t}{c_t} - \frac{\sigma_{cq,t}}{q_t(1+\tau_t^q)} \right]$$

Define $\mathcal{I}_{c,t}^q = \frac{\partial}{\partial c_t} \ln\left(\frac{U_{q,t}}{U_{c,t}}\right)$ and $\mathcal{I}_{y,t}^q = \frac{\partial}{\partial y_t} \ln\left(\frac{U_{q,t}}{U_{c,t}}\right)$, representing how the willingness-to-pay for the dirty good responds to changes in clean good consumption and in labor income, respectively. Note that, using the elasticity definitions from the main text,

$$\begin{aligned}
\frac{\partial}{\partial c_t} \ln\left(\frac{U_{q,t}}{U_{c,t}}\right) &= \frac{U_{cq,t}}{U_{q,t}} - \frac{U_{cc,t}}{U_{c,t}} = \frac{\sigma_t}{c_t} - \frac{\sigma_{cq,t}}{q_t} \frac{U_{c,t}}{U_{q,t}} = \frac{\sigma_t}{c_t} - \frac{\sigma_{cq,t}}{q_t(1+\tau_t^q)} \\
\frac{\partial}{\partial y_t} \ln\left(\frac{U_{q,t}}{U_{c,t}}\right) &= \frac{U_{qy,t}}{U_{q,t}} - \frac{U_{cy,t}}{U_{c,t}} = \frac{1}{\theta}\left(\frac{U_{ql,t}}{U_{q,t}} - \frac{U_{cl,t}}{U_{c,t}}\right) = -\frac{\gamma_{c,t} - \gamma_{q,t}}{y_t}
\end{aligned}$$

Defining $\zeta_t = (1 - \hat{\tau}_t^y)\eta_t - \eta_t^s$ and using the above equations,

$$\frac{\tau_t^q - \nu_t}{1+\tau_t^q} = (\gamma_{c,t} - \gamma_{q,t}) \frac{\psi U_{c,t}}{\theta f_t} - \eta_t \mathcal{I}_{y,t}^q - \zeta_t \mathcal{I}_{c,t}^q$$

$\frac{\psi U_{c,t}}{\theta f_t}$ is proportional to the labor wedge the planner would normally want to set. Hence, using the labor wedge derivations, we can also write

$$\frac{\tau_t^q - \nu_t}{1+\tau_t^q} = \frac{\gamma_{c,t} - \gamma_{q,t}}{1 + \epsilon_t - \gamma_{c,t}} \frac{\psi U_{c,t}}{\theta f_t} - \eta_t \mathcal{I}_{y,t}^q - \zeta_t \mathcal{I}_{c,t}^q$$

The only remaining thing to calculate is μ_t. Using $U_{Sl} l = \xi U_S$ and $U_{cS}/U_c = \xi_c (U_S/U_c)/c$,

$$-\frac{U_S}{U_c} \left[1 + (\xi - \gamma_c) \frac{\psi U_c}{\theta f} + \eta \frac{\xi - \gamma_c}{y} + ((1 - \bar{\tau}_t^y)\eta - \eta_t^s) \frac{\sigma + \xi_c}{c} \right]$$

Define $\mathcal{I}^S_{c,t} = \frac{\partial}{\partial S_t} \ln\left(\frac{U_{q,t}}{U_{c,t}}\right)$ and $\mathcal{I}^S_{y,t} = \frac{\partial}{\partial S_t} \ln\left(\frac{U_{q,t}}{U_{c,t}}\right)$, representing how the willingness-to-pay for environmental quality responds to changes in clean good consumption and in labor income. From here,

$$\begin{aligned}\mu_t &= -\mathbb{E}_0\left[\left(1 + (\xi - \gamma_c)\frac{\psi U_c}{\theta f} + \eta \mathcal{I}^S_y + \zeta_t \mathcal{I}^S_c\right)\frac{U_S}{U_c}\right] \\ &= -\mathbb{E}_0\left[\left(1 + \frac{\xi - \gamma_c}{1 + \epsilon - \gamma_c}\frac{\tau^y}{1-\tau^y} + \eta \mathcal{I}^S_y + \zeta_t \mathcal{I}^S_c\right)\frac{U_S}{U_c}\right]\end{aligned} \tag{20}$$

B.2 Savings multiplier over the life cycle

Fix a realized history. Let $\kappa_t = (1 - \bar{\tau}^s_t)\beta R^2 > 0$, $P_t = \prod_{j<t} \kappa_j$, with $P_0 = 1$, and $x_t = \vartheta_t / P_t$. Equation (12) gives

$$x_{t+1} = x_t - \frac{\eta^s_t}{P_t U_{c,t}}, \qquad t < T. \tag{21}$$

The endpoint conditions $\vartheta_0 = \vartheta_{T+1} = 0$ imply $x_0 = 0$ and $x_T = \eta^s_T/(P_T U_{c,T})$. Summing (21),

$$\sum_{t=0}^{T} \frac{\eta^s_t}{P_t U_{c,t}} = 0$$

If any multiplier along the history is nonzero, terms of both signs must appear. We have

$$\frac{1 - \chi_t}{U_{c,t}} = \frac{1}{\beta R}\mathbb{E}_t\left[\frac{1 - \chi_{t+1}}{U_{c,t+1}}\right] \tag{22}$$

If we define $\mathcal{G}_\sqcup \equiv 1 - \gamma_{c,t}\frac{\psi U_{c,t}}{\theta f_t} + \eta_t \Gamma_{c,t}$,

$$1 - \chi_t = \mathcal{G}_t - \eta^s_t \frac{\sigma_t}{c_t} \tag{23}$$

Where $\mathcal{G}_t > 0$ and $\mathbb{E}_t[\mathcal{G}_{t+1}] > 0$, define the reference wedge by $1 - \tau^{s,*}_t = \mathcal{G}_t / \mathbb{E}_t[\mathcal{G}_{t+1}]$. Multiplying (22) by $U_{c,t}$ and substituting the Euler equation with the fixed savings wedge gives

$$1 - \chi_t = (1 - \bar{\tau}^s_t)\left\{\mathbb{E}_t[1 - \chi_{t+1}] - \mathrm{Cov}_t\left(U_{c,t+1}, \frac{1 - \chi_{t+1}}{U_{c,t+1}}\right)\right\}$$

Substitution of (23) then yields

$$\eta_t^s \frac{\sigma_t}{c_t} - (1-\bar{\tau}_t^s)\,\mathbb{E}_t\left[\eta_{t+1}^s \frac{\sigma_{t+1}}{c_{t+1}}\right] = \mathbb{E}_t[\mathcal{G}_{t+1}](\bar{\tau}_t^s - \tau_t^{s,*}) + (1-\bar{\tau}_t^s)\,\mathrm{Cov}_t\left(U_{c,t+1}, \frac{1-\chi_{t+1}}{U_{c,t+1}}\right)$$

When the next draw is degenerate this reduces to

$$\eta_t^s \frac{\sigma_t}{c_t} - (1-\bar{\tau}_t^s)\eta_{t+1}^s \frac{\sigma_{t+1}}{c_{t+1}} = \mathcal{G}_{t+1}(\bar{\tau}_t^s - \tau_t^{s,*}). \tag{24}$$

For a deterministic history, let $z_t = \eta_t^s \sigma_t / c_t$ and assume $\sigma_t > 0$, $\mathcal{G}_t > 0$, and $1 - \bar{\tau}_t^s > 0$. Strict curvature is needed because Assumption 2.2 alone allows $\sigma_t = 0$. Equation (24) reads

$$z_t = 1 - \bar{\tau}_t^s z_{t+1} \mathcal{G}_{t+1}(\bar{\tau}_t^s - \tau_t^{s,*})$$

If every $g_t > 0$, then $z_t \leq 0$ implies that $z_{t+1} < 0$. We know some multipliers mut be nonzero, so the sequence starts positive and ends negative, crossing once. If every $g_t < 0$, this reverses.

C Derivations: linear carbon price

C.1 Expenditure, demand, and incentive compatibility

To simplify notation, let $\bar{p}_t \equiv 1 + \bar{\tau}_t^q$. The planner has the below dirty good wedge constraint:

$$\Gamma_t^q \equiv \frac{U_{q,t}}{U_{c,t}} - \bar{p}_t = 0$$

At fixed price, labor, and pollution, define indirect utility by

$$\tilde{U}(x, l; \bar{p}, S) = \max_{c,q} \; U(c, q, l, S) \quad \text{s.t.} \quad c + \bar{p}q = x$$

Here, x is the total expenditure on the two goods. If we substitute $c = x - \bar{p}q$, we get the first-order condition $U_q = \bar{p}U_c$. With a linear dirty good wedge, the planner is essentially only choosing each household's total expenditure x and labor l rather than also the consumption mix, which is now equal for everyone. If we write $\tilde{U}(x, y/\theta; \bar{p}_t, S_t)$, we have $\tilde{U}_x = U_c$ because we can express x as a function of c. Now that the planner cannot impose a consumption mix on households, it chooses the linear dirty good wedge and total expenditure such that incentive constraints are satisfied.

$$\tilde{U}\left(x(\hat{\theta}), \frac{y(\hat{\theta})}{\theta}; \bar{p}_t, S_t\right) \geq U\left(c(\hat{\theta}), q(\hat{\theta}), \frac{y(\hat{\theta})}{\theta}, S_t\right)$$

Differentiate $U_q = \bar{p}U_c$ jointly with the expenditure constraint. Define $\mathcal{D} \equiv U_{qq} - 2\bar{p}U_{cq} + \bar{p}^2 U_{cc}$. At fixed labor and pollution, expenditure and compensated price changes give

$$\begin{aligned} \frac{\partial q^d}{\partial x} &= \frac{U_{cq} - \bar{p}U_{cc}}{|\mathcal{D}|} = \frac{\bar{p}U_c \mathcal{I}_c^q}{|\mathcal{D}|} \\ s_{qq} \equiv \left.\frac{\partial q^d}{\partial \bar{p}}\right|_u &= \frac{U_c}{\mathcal{D}} < 0 \end{aligned}$$

From here, we can write

$$\Delta_t^q \equiv \bar{p}\mathcal{I}_c^q - \mathcal{I}_q^q = -\frac{\mathcal{D}}{\bar{p}U_c} = \frac{1}{\bar{p}|s_{qq}|} > 0 \tag{25}$$

C.2 Optimal linear dirty good wedge

The Hamiltonian of Appendix A.1 gains $\eta^q f \Gamma^q$. Using the usual willingness-to-pay identities,

$$\Gamma^q_z = \frac{U_{qz} - \bar{p} U_{cz}}{U_c} = \bar{p} \mathcal{I}^q_z. \tag{26}$$

Each pointwise condition for c, q, and l gains $\eta^q f \Gamma^q_z$. With income wedges free,

$$\frac{\phi U_c}{f} = 1 - \gamma_c \frac{\psi U_c}{\theta f} + \eta^q \bar{p} \mathcal{I}^q_c, \tag{27}$$

$$\frac{\phi U_c}{f} = \frac{1 + \nu_t}{\bar{p}} - \gamma_q \frac{\psi U_c}{\theta f} + \eta^q \mathcal{I}^q_q. \tag{28}$$

Here ν_t is the shadow emission price in (13). Equating (27) and (28) gives

$$\underbrace{1 - \frac{1 + \nu_t}{\bar{p}_t}}_{\varpi_t} = \underbrace{(\gamma_c - \gamma_q) \frac{\psi U_c}{\theta f}}_{\varpi^*_t} - \eta^q_t \Delta^q_t.$$

Using (25),

$$\eta^q_t = \bar{p}_t |s_{qq,t}| (\varpi^*_t - \varpi_t).$$

The common wedge enters the date-0 resource-cost Lagrangian only through the date-t price constraints. Accounting for resource discounting and transition probabilities gives

$$\frac{d\mathcal{L}_0}{d\bar{\tau}^q_t} = -\frac{1}{R^t} \mathbb{E}_0[\eta^q_t].$$

The planner minimizes resource cost at fixed promised welfare, so a negative derivative favors raising the dirty good wedge. An interior optimum sets this derivative to zero. Since $\bar{p}_t$ and ϖ_t are common across date-t histories,

$$\varpi_t = \frac{\mathbb{E}_0[|s_{qq,t}| \varpi^*_t]}{\mathbb{E}_0[|s_{qq,t}|]}$$

Weak separability gives $\varpi^*_t = 0$ at every history, so the optimal date-specific rate is ν_t.

C.3 Fixed income wedges and a lifetime rate

The conditions for c, q, and l gain $\eta^q f \Gamma_z^q$. Equating the resulting conditions for c and q,

$$\varpi_t = \varpi_t^* - \eta_t \left(\Gamma_c - \frac{\Gamma_q}{\bar{p}_t} \right) + \eta_t^s \mathcal{I}_{c,t}^q - \eta_t^q \Delta_t^q.$$

Using the Γ definitions and rewriting as in Appendix ,

$$\Gamma_c - \frac{\Gamma_q}{\bar{p}_t} = (1 - \bar{\tau}_t^y) \mathcal{I}_{c,t}^q \frac{\gamma_{c,t} - \gamma_{q,t}}{y_t}$$

Putting the $\mathcal{I}_{c,t}^q$ terms together,

$$\varpi_t = \varpi_t^* + \eta_t \frac{\gamma_{c,t} - \gamma_{q,t}}{y_t} - ((1 - \bar{\tau}_t^y)\eta_t - \eta_t^s) \mathcal{I}_{c,t}^q - \eta_t^q \Delta_t^q$$

Alternatively,

$$\varpi_t = \varpi_t^* + \eta_t \frac{\gamma_{c,t} - \gamma_{q,t}}{y_t} - \zeta_t \mathcal{I}_{c,t}^q - \eta_t^q \Delta_t^q$$

With labor and savings wedges fixed, the linear dirty good wedge satisfies

$$\varpi_t = \frac{\mathbb{E}_0[|s_{qq,t}| \ \varpi_t^{**}]}{\mathbb{E}_0[|s_{qq,t}|]}$$

If the planner must not only hold the dirty good wedge constant across households in a given date but also across dates, summing C.2 across dates and substituting the multiplier formula,

$$\sum_{t=0}^{T} \frac{1}{R^t} \mathbb{E}_0[|s_{qq,t}|(\varpi_t^{**} - \varpi_t)] = 0$$

For a time-varying social cost of carbon, $\varpi_t = (\bar{\tau}^q - \nu_t)/\bar{p}$ and this is equivalent to

$$\bar{\tau}^q = \frac{\sum_{t=0}^{T} \frac{1}{R^t} \mathbb{E}_0[|s_{qq,t}|(\nu_t + \bar{p}\varpi_t^{**})]}{\sum_{t=0}^{T} \frac{1}{R^t} \mathbb{E}_0[|s_{qq,t}|]}.$$

C.4 Social cost of carbon

The linearity constraint affects the damage weight through both the direct pollution derivative and the delivery cost ϕ. From (26) and (27), their combined contribution at each history is

$$\eta^q \left(\Gamma_S^q - \Gamma_c^q \frac{U_S}{U_c} \right) = -\eta^q \frac{U_S}{U_c} \left(\bar{p} \mathcal{I}_c^S - \mathcal{I}_q^S \right).$$

This follows directly from $\mathcal{I}_z^S = \partial_z \ln(-U_S/U_c)$ and the quotient rule. With income wedges free, substituting into the condition of Appendix A.5 yields

$$\mu_t = -\mathbb{E}_0 \left[\left(1 + (\xi - \gamma_c) \frac{\psi U_c}{\theta f} + \eta^q \left(\bar{p} \mathcal{I}_c^S - \mathcal{I}_q^S \right) \right) \frac{U_S}{U_c} \right].$$

The expectation averages all date-t histories. With income wedges fixed, this adds the same η^q term to the bracket of (20). The emission recursion (13) is unchanged, so aggregating these marginal stock damages as in Appendix A.5 gives the social cost of carbon.

Choosing the common wedge optimally does not generally remove the additional term. At a date whose price is free to adjust, $\mathbb{E}_0[\eta_t^q] = 0$ and its contribution to the social cost of carbon is

$$\mathbb{E}_0[\eta_t^q \mathcal{K}_t] = \mathrm{Cov}_0(\eta_t^q, \mathcal{K}_t) \qquad \mathcal{K}_t \equiv -\frac{U_{S,t}}{U_{c,t}} \left(\bar{p}_t \mathcal{I}_{c,t}^S - \mathcal{I}_{q,t}^S \right).$$

The social cost remains common across households because pollution is aggregate, but its level can change with the constrained allocation.